

Algorithms and Responsibility

Felix Bönisch*

July 6, 2026

Abstract

A long line of experimental work suggests that decision-makers can—and do—use delegation to shift responsibility for unwelcome outcomes onto an intermediary. We ask whether the same logic carries over to delegation to algorithms. In a preregistered online experiment ($N = 322$ UK Prolific participants, 161 in each role), a decision-maker makes a payoff-relevant prediction for a matched recipient with the option to delegate to an algorithm calibrated, at the individual level, to perform as well as she does. We manipulate, between subjects, whether the matched recipient can punish her for the decision. Two findings emerge. First, contrary to the preregistered prediction of the responsibility-avoidance literature, the *prospect* of punishment *reduces* rather than increases the likelihood of delegating to the algorithm (a 17 percentage-point gap, two-sided $p = 0.033$). Second, conditional on outcome, we cannot reject equal punishment of delegated and self-made decisions; the point estimate runs in the predicted direction but the test is underpowered to rule out moderate insulation effects. We discuss several candidate mechanisms for the headline reversal and argue that, among those we can evaluate, a process-ownership account—in which the moral salience introduced by the punishment possibility raises the value of being the proximate cause of a good outcome—is the most consistent with the data, while flagging that the supporting evidence is exploratory and the design cannot rule out experimenter demand. Within the bounds of the setting we study, our results offer a behavioral complement to regulatory frameworks that hold human users accountable for algorithm-mediated decisions: such frameworks may discipline algorithmic adoption rather than merely allocate post-hoc responsibility.

JEL Classification: C91, D81, D91, O33

Keywords: algorithms; delegation; responsibility; punishment; experiment

*WZB Berlin Social Science Center. Email: felix.boenisch@wzb.eu. Acknowledgements to be added at submission.

1 Introduction

Algorithms are increasingly involved in consequential decisions: they screen job candidates, set credit scores, recommend medical treatments, evaluate parole applicants, and—more mundanely but at vast scale—route the cars driving toward this very page. As the locus of decision-making shifts toward algorithmic systems, a parallel question has moved into public, legal, and academic debate: when an algorithm-mediated decision turns out badly, who bears the blame? A natural intuition, which finds support in a long line of experimental work on responsibility avoidance through delegation, is that decision-makers can use algorithms as a means to shift blame onto a non-human intermediary. To return to a familiar example: if you and a friend arrive late because she relied on a navigation app rather than her own sense of direction, the algorithm seems a convenient scapegoat—and the next time she is unsure of the route, the prospect of avoiding your blame may itself motivate her use of the app. This paper provides experimental evidence on whether either of these intuitions is borne out.

We address two questions. *First*: can decision-makers actually shift responsibility away from themselves by delegating to an algorithm—that is, are they punished less for a delegated decision than for a self-made decision with the same outcome? *Second*: does the prospect of being held responsible motivate the decision to delegate in the first place? The two questions correspond to, respectively, the *viability* and the *motivational salience* of algorithmic responsibility avoidance. They have been studied in different settings and rarely in combination, and previous designs have struggled to identify them cleanly. We do so here in a preregistered online experiment.

Our experimental design pins down both questions. Player A first completes ten rounds of an incentivised visual prediction task (estimating a person’s weight from a facial image and stated height). She then makes an final prediction whose outcome determines the payoff of a second participant, Player B: a high payoff if the prediction falls within ten pounds of the true weight, a low payoff otherwise. Player A may either make this prediction herself or delegate it to an algorithm. The algorithm is calibrated to mimic Player A’s individual performance from the first ten rounds—that is, on the final prediction it produces the high outcome with the same probability as Player A herself. Player A is told this. We manipulate the punishment possibility between subjects: in the *Punishment* condition, Player B may reduce Player A’s payoff conditional on the delegation decision and the realized outcome; in the *No-Punishment* condition, no punishment is possible. Player B’s punishment behavior is elicited via the strategy method across all four (delegation, outcome) cells, providing a within-subject test of whether delegation insulates Player A from punishment.

Three design features distinguish our setting from the closest precedent, [Feier et al. \(2022\)](#), two of them central to identification. First, the algorithm is calibrated at the *individual* level rather than the session or group level. Existing designs typically expose subjects to a distribution of algorithmic performance and let them infer their own ability relative to that distribution; the delegation decision then carries strong signal value about the principal’s self-perception,

and any differential punishment can reflect overconfidence rather than responsibility shifting per se. By neutralizing relative-performance signaling at the individual level, we isolate the responsibility-shifting motive. Second, our *No-Punishment* condition turns off the punishment possibility entirely, allowing us to identify the effect of anticipated punishment on the delegation decision itself—a test that none of the existing algorithmic-delegation designs permits. Third, we separate the delegator and evaluator roles across disjoint subject pools, rather than having the same subjects play both, removing the self-justification motives that a dual role can introduce into the punishment decision. Together these features sharpen the test of responsibility shifting in algorithmic delegation that the literature has thus far been able to deliver.

Two findings emerge, and in their joint structure they push against the natural reading of the responsibility-avoidance literature. *First, the prospect of punishment reduces, rather than increases, the likelihood of delegation:* 42.5% of Player As delegate when punishment is possible, compared to 59.3% when it is not (a 17 percentage-point gap; two-sided Pearson Chi-squared $p = 0.033$, Fisher exact $p = 0.041$; the preregistered one-sided test in the H1 direction is non-significant by construction, since the result reverses the predicted direction—we discuss this in detail in Section 4). This rejection of the standard prediction in favor of the opposite direction is the paper’s main empirical contribution, and we view it as a behavioral counterpart to the regulatory intuition that holding humans accountable for algorithm-mediated decisions disciplines their use of those algorithms. *Second, conditional on outcome, we cannot reject equal punishment of delegated and self-made decisions in our setting:* Player B’s punishment of Player A is statistically indistinguishable across the two, and Player A correctly anticipates this absence of differential punishment in her elicited beliefs. The point estimate of any insulation effect runs in the predicted direction but the test is underpowered against moderate alternatives ($n \approx 67$ Player Bs after attention-check exclusion); the broader claim that responsibility shifting through algorithmic delegation is altogether absent is one we cannot make with this data.

What mechanism reconciles the two findings? Among the candidates we can evaluate with our design, a *process-ownership* account is the most consistent with the data—though we treat it as a candidate hypothesis rather than an established mechanism, since the supporting evidence is exploratory and an experimenter-demand account cannot be ruled out by construction. A simple anticipation story—in which Player A delegates less when punishment is on the table because she expects harsher punishment for delegated decisions—is inconsistent with the elicited beliefs and with Player B’s actual choices. A simple effort story—in which Player A makes the prediction herself in order to outperform the algorithm and earn the high payoff—is inconsistent with the data, in which non-delegators in the *No-Punishment* condition outperform non-delegators in the *Punishment*. What remains is the milder hypothesis that, when the moral salience of the decision is heightened by the punishment possibility, Player A wants to be the proximate cause of any good outcome for Player B, irrespective of any quantifiable belief about how she will be judged. Consistent with this interpretation, the negative correlation between Player A’s task performance and her propensity to delegate is observed in the *Punishment* condition only: those most likely to deliver the high payoff want to do so personally when the

moral stakes are made salient. Section 4.1 weighs the four candidate mechanisms in detail and is explicit about what the data can and cannot rule out.

These findings contribute to three strands of literature. First, we add to the experimental literature on responsibility avoidance through delegation (Bartling and Fischbacher, 2012; Coffman, 2011; Oexl and Grossman, 2013; Hamman et al., 2010; Hill, 2015; Steffel et al., 2016), by extending the canonical setting to delegation to an algorithmic intermediary in a setting with genuine outcome uncertainty. Second, we contribute to the rapidly growing literature on algorithm aversion and adoption (Dietvorst et al., 2015; Logg et al., 2019; Burton et al., 2020; Mahmud et al., 2022; Chugunova and Sele, 2022) by isolating an under-studied determinant of algorithm use: the institutional context in which the human user is held accountable. Third, and most directly, we build on Feier et al. (2022) and Kirchkamp and Strobel (2019) by providing a clean experimental test of two long-standing claims in the responsibility-and-algorithms literature, with a design that controls relative-performance signaling at the individual level and that includes a condition in which punishment is ruled out by construction. The contrast with Feier et al. (2022) is consequential: that study finds that decision-makers are held *less* responsible for algorithmically delegated decisions, while we find that the punishment possibility *reduces* delegation rather than increasing it. We argue that the two findings are reconcilable: the relative-performance signaling that drives much of Feier et al.’s differential punishment is closed off by construction in our design, isolating a different and milder margin (the motivational salience of accountability for the principal). Beyond the experimental literature, our findings speak to the regulatory debate over the assignment of responsibility for algorithmic decisions (European Union, 2024): within the bounds of the setting we study, holding human users accountable for algorithm-mediated outcomes appears to be a behavioral lever for direct human engagement with the decision, rather than merely a tool for ex-post moral attribution. We are explicit in Section 5 about why the present evidence—a single online experiment on a stylised prediction task with one UK Prolific sample—supports this claim only in a bounded sense, and what would be required to extend it to the high-stakes algorithmic domains (medical, judicial, hiring, financial) that motivate the regulatory debate.

The remainder of the paper is organized as follows. Section 2 reviews the relevant literature in detail. Section 3 describes the experimental design and states our two preregistered hypotheses. Section 4 presents the empirical findings, including the mechanism analysis. Section 5 discusses the implications and limitations of our results.

[FB: Include a footnote on pre-registration.]

2 Related Literature

Responsibility avoidance through delegation

Delegation has long been studied as an efficiency-enhancing institutional arrangement: principals delegate to agents who have better information, lower opportunity costs, or specialised skills [Comment: From Bartling and Fischbacher "for an overview of the principal-agent literature see, e.g. Bolton and Dewatripont, 2005"]. Alongside these efficiency motives, the literature has identified a range of strategic motives for delegation, e.g. as a commitment device that alters the equilibrium (Schelling, 1960; Fershtman and Gneezy, 2001) [Comment: Bartling and Fischbacher cite Schelling (1960) here. In a footnote, they then specify: "Applications are output and pricing decisions in oligopolistic markets (Vickers, 1985), inflation targeting (Rogoff, 1985), and bargaining (Jones, 1989). Huck, Müller and Normann (2004) provide experimental evidence on strategic delegation in oligopolistic markets, Schotter, Zheng and Snyder (2000) and Fershtman and Gneezy (2001) in bargaining situations"], as a way to motivate the agent (Aghion and Tirole, 1997), or as a means of shifting responsibility for the outcome away from the principal, i.e. principals may delegate not because the agent is more competent or cheaper, but because doing so shifts responsibility for the resulting outcome away from themselves — both in their own eyes (felt responsibility for the outcome) and in the eyes of others (the blame others attribute to them). The political-science and public-choice literatures have long recognized that delegation can serve as a strategy for shifting blame onto an intermediary (Weaver, 1986; Fiorina, 1986). Experimental economics has established the same motive in incentivised settings, primarily using modified dictator games.

Bartling and Fischbacher (2012) provide the canonical experimental evidence. In their design, a dictator may either choose between a fair and an unfair allocation of an endowment or instead delegate this choice to another person whose monetary incentives are aligned with the dictator's. A recipient, who is adversely affected by the unfair allocation, may then assign costly punishment to the dictator and the delegee. They find that through delegation punishment is effectively shifted from the dictator to the intermediary, and the prospect of avoiding punishment constitutes a strong motive for the delegation decision itself. Coffman (2011) reinterprets the mechanism behind this punishment reduction. He finds that the principal can no longer avoid punishment when she directly interacts with the receiver despite the intermediary's presence, suggesting that what reduces punishment is the absence of direct interaction with the affected party rather than any transfer of blame to a deciding intermediary. This holds even when the intermediary remains completely passive. Oexl and Grossman (2013) sharpen this reading by modifying the canonical game so that delegation necessarily eliminates the fair option. The principal nevertheless escapes punishment, with the intermediary actively choosing between two unfair allocations after delegation. Finally, both Bartling and Fischbacher (2012) (with a die roll) and Oexl and Grossman (2013) (with an intermediary-triggered random draw) include treatments that substitute a randomization device for the human intermediary. In both, the

principal still shifts some punishment to the device, though less than under delegation to a deciding human. Consistent with this, in [Bartling and Fischbacher \(2012\)](#) fewer dictators delegate to the die than delegate to a human delegatee in the corresponding treatment. Two components of responsibility shifting thus emerge — the absence of direct interaction between principal and recipient drives the reduction in principal punishment, while the act of choosing, rather than any contribution to outcome unfairness, is what transfers blame to the intermediary.

The motive is not confined to settings with a punishment stage. [Hamman et al. \(2010\)](#) extend this picture in a setting without punishment opportunity. Still, principals’ transfers to recipients fall sharply when they delegate, indicating that responsibility-shifting can operate through felt responsibility or moral wiggle room rather than through the avoidance of punishment. Qualitative self-reported responsibility and recipient-side blame attributions are consistent with this reading—dictators feel less responsible when they delegate, and recipients place less blame on principals than on intermediaries.

Beyond the canonical dictator-game setting, the responsibility-shifting motive — and the strategic use of delegation it gives rise to — has been documented across a range of domains. In incentivised ultimatum bargaining, proposers using a hired delegate extract larger surpluses ([Fershtman and Gneezy, 2001](#)). In sequential voting, non-pivotal voters are punished less than pivotal ones for the same group outcome and voters strategically vote against their preferred allocation to avoid being pivotal ([Bartling et al., 2015](#)). For acts of deception, principals pay an agent to lie on their behalf ([Erat, 2013](#)) [Comment: see also Kandul and Kirchkamp, 2018; Gawn and Innes, 2019a; maybe also include communication bad information/news ([Garofalo and Rott, 2018](#)); workplace ([Charness \(2004\)](#) and [Maximiano et al. \(2013\)](#))]. In policy-making, observers rate delegating legislatures as less blameworthy than legislatures who decide directly, even when the delegate is powerless to change the outcome ([Hill, 2015](#)). [Freer et al. \(2024\)](#) isolate the motive in financial decision-making, finding that a substantial fraction of investors delegate even purely luck-based lottery choices, where alternative motives — decision costs, performance-chasing, or risk tolerance — cannot rationalize the choice.

[Comment: - "For example, individuals prefer to manage their own stock portfolio instead of delegating the decisions to a financial advisor, despite making frequent emotional trades and lower returns ([Barber and Odean, 2001](#))" - "While some studies show that individuals are less likely to delegate to AI than to humans (e.g., [Dietvorst et al., 2015](#); [Leyer and Schneider, 2019](#)), others find a preference for AI over human agents in strategic decision settings (e.g., [Logg et al., 2019](#); [Candrian and Scherer, 2022](#); [Germann and Merkle, 2023](#); [Holzmeister et al., 2023](#))."]

In sum, the canonical literature documents that in settings with self-interested choice, delegation reduces both the principal’s punishment and her felt responsibility, in turn motivating delegation. Our setting departs from this canonical design along three dimensions: the intermediary is algorithmic rather than human, the outcome falls entirely on a passive recipient rather than partly on the principal herself, and the underlying decision involves genuine predictive uncertainty rather than a known fair–unfair choice. Whether the responsibility-shifting motive

retains its force under these conditions is the question we take up.

Delegation to algorithms

[Comment: Perhaps distinguish more strictly between delegation vs. advice-taking. Or include it into the "ceding control" subsegment]

A large and rapidly growing empirical literature studies when and why people are willing to use algorithmic decision aids. Dietvorst et al. (2015) introduce the term *algorithm aversion* — the finding that people prefer their own (or other humans') judgement over algorithmic predictions, even after observing that algorithms outperform human decision-makers.¹ In contrast, Logg et al. (2019) document conditions under which subjects show *appreciation* of algorithmic advice over human advice. The picture that emerges is therefore one of substantial heterogeneity in algorithm preferences across tasks, contexts, and elicitation methods. Burton et al. (2020), Mahmud et al. (2022), Chugunova and Sele (2022), and Qin et al. (2025) provide systematic reviews of this literature and document a large array of moderators for when and why algorithms are used in decision-making.

For example, algorithm aversion is stronger when the algorithm cannot be observed to learn from its mistakes (Berger et al., 2021), when the algorithm's reasoning is opaque (Litterscheidt and Streich, 2020; Sharan and Romano, 2020), when the algorithm is presented as more complex (Stein et al., 2020), when the task is perceived as more uncertain (Dietvorst and Bharti, 2020), when delegation removes any ability to override the algorithm's choice (Dietvorst et al., 2018), when the algorithm is used in tasks perceived as subjective (Castelo et al., 2019; Bogert et al., 2021), and when decision tasks are framed as a loss (Bockstedt and Buckman, 2025).

For example, on the algorithm side, algorithm aversion is stronger when the algorithm cannot be observed to learn from its mistakes (Berger et al., 2021) and when its reasoning is opaque (Litterscheidt and Streich, 2020; Sharan and Romano, 2020).

On the task side, algorithm aversion rises with the degree of irreducible uncertainty (Dietvorst and Bharti, 2020), with the perceived subjectivity of the task (Castelo et al., 2019), and with loss framing of the decision (Bockstedt and Buckman, 2025), while reliance on algorithmic advice increases with task difficulty (Bogert et al., 2021).

On the control side, people use imperfect algorithms far more readily when they retain even a minimal ability to modify the output (Dietvorst et al., 2018) [<https://www.sciencedirect.com/science/article/pii/S0749597823000262>] [normann_delegate₂025]

[Include soewhere: <https://www.sciencedirect.com/science/article/abs/pii/S0749597823000262>]

[Comment: - the literature reviews and references in general are old-ish...ideally we can include

¹Interest in the comparison between human and statistical judgement is long-standing — see Meehl (1954) and Dawes (1979) for the classic demonstrations that simple statistical rules outperform clinical judgement and are resisted nonetheless.

more recent references. - Maybe include the Dargies, Hakimov, Kübler reference here. - Also think about where to include the two ivanova stenzel papers]

Responsibility perceptions of decisions involving algorithms

[<https://www.annualreviews.org/docserver/fulltext/psych/75/1/annurev-psych-030123-113559.pdf?expire>

Within this body of work, a smaller subset focuses on algorithm use preferences in settings whose consequences fall partly or wholly on someone other than the decision-maker, sometimes called the *moral domain* (Gogoll and Uhl, 2018). Much of this work concludes that aversion to algorithms is amplified once third-party welfare is at stake. Gogoll and Uhl (2018) document this in an incentivised calculation task with performance uncertainty (both humans and the algorithm can err), Bigman and Gray (2018) in vignettes across high-stakes moral domains (driving, legal, medical, military), and Jauernig et al. (2022) identify *moral discretion* as the human capacity that people specifically value in such settings.²

Once decisions affect a third party, the responsibility-shifting motive established for human delegation becomes relevant for algorithmic delegation as well — and with it the question of how observers apportion blame when an algorithm is involved in the decision. A first strand of evidence, predominantly vignette-based and unincentivised, examines how third parties attribute responsibility and blame to human actors involved with algorithms. In the medical domain, Pezzo and Pezzo (2006) document that a physician who follows a (bad) recommendation from a computerised decision aid is faulted less than one who errs unaided. Tobia et al. (2021) extend the discussion to liability law: physicians who follow standard-care AI recommendations face reduced legal exposure — and the authors argue that this constitutes an incentive to follow such recommendations specifically because doing so shields physicians from blame. In hiring, algorithmic discrimination triggers less moral outrage than identical human discrimination, and the responsible organization is judged less harshly (Bigman et al., 2023). In the workplace setting, Naquin and Kurtzberg (2004) find that observers attribute less responsibility to a company when an accident involves technology. [Comment: This is rather old. I would like to include references only after 2020 for the individual domains] Across multiple moral domains (driving, legal, medical, military), Bigman and Gray (2018) document a robust “algorithmic outrage asymmetry”: moral violations by algorithms produce less outrage than identical violations by humans. Shank et al. (2019) consider moral violations produced by human-algorithm pairs and find that humans who monitor algorithms are blamed less than humans acting alone or humans receiving algorithm recommendations, suggesting that responsibility attributions are sensitive to the precise division of authority between human and machine.

Evidence using incentivised laboratory experiments is younger and smaller. Kirchkamp and Strobel (2019) compare a modified dictator game in which the dictator either acts alone or

²For recent surveys of this area see Bonnefon et al. (2024) and, with a legal-policy focus, Sunstein and Gaffe (2024).

makes the decision jointly with a computer. They find no significant difference in the dictator’s self-felt responsibility between human-computer and human-human teams, although a tendency towards more selfish allocations under shared decision-making with humans. [Maasland and Weißmüller \(2022\)](#) find in a computer-aided HR setting that algorithm aversion *dominates* blame avoidance: the prospect of using an algorithm to delegate an unpleasant HR decision (a dismissal) does not overcome the underlying aversion to algorithms in this domain.

[Comment: It’s hard to follow...there is no real red line here. The first reference is about comparing human-human to human-algorithm, which is yet another dimension. We are not talking about joint decisions, but full delegation. Nevertheless I think it’s nice to include it. Does it not tie into Shank attributions 2019?]

A more recent strand documents the opposite pattern: algorithmic intermediaries serving as moral cover. [Chevrier and Teixeira \(2024\)](#) extend the blame chain to include the programmer behind the algorithm: in their design, AI programmers and the delegator both escape punishment for inequalitarian AI-mediated allocations, and the resulting “moral wiggle room” is exploited. [Comment: ”...extend the blame chain to the programmer behind the algorithm: recipients punish inequalitarian AI-mediated allocations less often, the programmer of an autonomous AI is punished less severely than the programmer of a fully controlled one — and programmers respond by choosing the inequalitarian option more often, exploiting the moral wiggle room that the AI’s autonomy creates, while delegation rates themselves do not differ across intermediary types.”] [Tontrup and Sprigman \(2025\)](#) examine delegation of a dictator-game allocation to ChatGPT and find that delegators transfer less to the recipient than non-delegators, indicating that AI delegation is used strategically to extract surplus while shedding moral responsibility, similarly to delegation to other humans in the canonical literature. [Hüholt and Szech \(2026\)](#) use an incentivised trolley-style donation choice between two real charitable allocations and find that subjects delegate the moral decision to an AI significantly more often than to a human counterpart. Across these four different settings, the AI-as-blame-shield motive is shared, and the direction of the effect on delegation is uniformly positive — in contrast to [Maasland and Weißmüller \(2022\)](#), where algorithm aversion dominates the responsibility-shifting motive in the same broad domain.

The most closely related study is [Feier et al. \(2022\)](#), who study delegation to an algorithm in an incentivised laboratory experiment. In their design, to determine a third party’s payoff subjects have the option to either rely on their own performance in a Raven-style logic puzzle or, depending on the treatment, delegate to another human or an algorithm. The subject’s own payoff is unaffected by this choice. Third parties may then reward or punish the principal. The headline finding is that decision-makers are held *less* responsible for decisions delegated to an algorithm than for decisions delegated to another human, conditional on outcome. Interestingly, in their experiment, punishment is only reduced significantly when delegating to an algorithm, but not to another human, compared to self-made decisions. Thus, algorithm-delegation, but not human-delegation, shields the principal from blame for bad outcomes.

Our design extends [Feier et al. \(2022\)](#) along three dimensions. First, where their algorithm is calibrated at the session level so that the delegation decision reflects beliefs about own ability *relative* to that distribution — a signal value about Player A’s self-perception that may contaminate downstream punishment as a measure of responsibility shifting per se — we calibrate the algorithm at the individual level, removing the relative-performance signal of delegation by construction. [Comment: ”First, their algorithm is calibrated at the session level, so the delegation decision carries a signal about the subject’s beliefs concerning her ability relative to the pool — indeed, their own analysis shows that the propensity to delegate is driven primarily by self-assessed errors in the logic task, not by the type of delegate. Punishment of a non-delegator may then partly sanction perceived overconfidence rather than responsibility-taking per se. We instead calibrate the algorithm to each subject’s own measured performance and make this common knowledge between the players, removing the relative-performance signal of delegation by construction.”] Second, where the same subject acts as both delegator and evaluator in their design — potentially introducing self-justification, hypocrisy aversion, and similar motives into the punishment decision — we separate the delegator and evaluator roles across disjoint subject pools.³ Third, their design does not include a treatment in which punishment is ruled out by construction, so that delegation rates between the two treatments may be different not because of differences in anticipated punishment, but more generally by preferences for algorithms. Thus, the motivational question of whether the prospect of punishment itself drives delegation is not identified. We add this manipulation directly, isolating the role of anticipated punishment in the delegation decision itself.

Our experiment is, to our knowledge, the first to our knowledge to provide a clean test of blame-shifting onto an algorithm in a setting with inherent uncertainty, controlling for relative-performance concerns. It is also the first to vary the punishment possibility for an algorithmic delegate, and hence to test whether, in such a setting the prospect of punishment itself motivates the delegation decision.

[Comment: Include somewhere about Feier: - The punishment gap is not specifically outcome-conditional — observers reward AI use in good outcomes as well as bad. This is consistent with a decision-rewarding rather than a blame-shielding reading. - Regarding the double role of subjects: delegation decision is significantly correlated with punishment difference for good outcomes.]



3 Experimental Design

We ran an online experiment via the British platform Prolific (<https://www.prolific.com>) with UK participants. The experiment was programmed in oTree ([Chen et al., 2016](#)) and eligibility followed Prolific’s standard quality screens as well as a minimum 95% prior approval rate, a

³For example, in their data, the evaluator’s own delegation decision is significantly correlated with the evaluation gap between delegated and non-delegated decisions after good outcomes ($r = -0.34$ in the machine treatment); own computations from their published data)

fluent-English filter and a desktop-only restriction (no mobile or tablet) to keep the screen layout consistent across subjects.

The design combines a between-subject manipulation of whether costly punishment is possible with a within-subject elicitation of conditional punishment behavior. Within each condition, participants assumed one of two roles. Player A decides whether to make a payoff-relevant prediction for Player B directly, or to delegate the prediction to an algorithm calibrated to match her own measured performance. Player B is the matched recipient whose payoff is fully determined by the outcome of Player A’s prediction, whether Player A makes it herself or delegates it to the algorithm. In the *Punishment* condition Player B may subsequently punish Player A, conditional on both the delegation decision and the realized outcome, while in the *No-Punishment* condition no punishment is possible. Figure 1 summarizes the design.



Procedure

Player A

After providing informed consent, Player A completed ten incentivized rounds of a visual prediction task. Specifically, in each round, she predicted a person’s weight from a facial image and the person’s height. Predictions could be entered via a slider or typed directly in kilograms,

pounds, or stones-and-pounds.⁴ No feedback was provided between rounds. At the end of the experiment, one of the ten rounds was randomly selected for payment: Player A earned £4 if her prediction in that round was within 10 lbs of the true weight and £2 otherwise. These initial rounds served three purposes simultaneously: First, they established a stake from which any later punishment could be deducted. Second, they familiarized Player A with the task. Third, they generated the performance data needed to calibrate the algorithm.

The choice of a visual prediction task is not incidental. Algorithm preferences may depend sensitively on the algorithm’s perceived relative competence (Qin et al., 2025; Logg et al., 2019). Our design neutralizes this concern by calibrating the algorithm to match each subject’s individual performance. The visual nature of the task lends credibility to this equal-performance claim, as people have everyday familiarity with judging others’ build and appearance, and algorithmic error is plausible.⁵ In contrast, subjects may be skeptical that an algorithm could ever fail in purely mathematical tasks (Gogoll and Uhl, 2018; Feier et al., 2022). Finally, we chose a prediction task involving genuine uncertainty rather than a setting with an obvious self-serving motive (as in classic dictator-game variations). This reflects the nature of many real-world applications of algorithms—financial forecasting, medical diagnosis, fraud detection—in which outcomes are inherently uncertain and performance is evaluated *ex post*.

After the ten rounds, Player A was randomly paired with another participant, Player B, who had not completed the prediction task himself. Player A was informed that her final prediction would not affect her own payoff but would almost entirely determine Player B’s. A prediction within 10 lbs of the true weight would give Player B a high payoff (£5), and any less accurate prediction a low payoff (£1). Player A then learned that she could either make the prediction herself or delegate it to an algorithm. She was also told that, based on her first ten predictions, the algorithm would have performed equally well.

The algorithm was calibrated to mimic Player A’s own performance over the first ten rounds. Concretely, if Player A delegates the decision to the algorithm, the algorithm was equivalent to drawing a Bernoulli outcome with success probability equal to Player A’s empirical accuracy rate over her first ten rounds. For example, a subject who scored 3 out of 10 in the initial rounds thus delegates to an algorithm that produces the high payoff for Player B with 30% probability.

This individual-level calibration is a central design choice. Fixing the algorithm’s success probability to each subject’s own accuracy makes the equal-performance claim hold subject by subject, not merely on average across a pool of participants. This matters because beliefs about relative performance are otherwise a dominant driver of delegation.⁶ For example, under the session-level calibration of Feier et al. (2022), the delegation choice is largely driven by

⁴All individuals depicted in the images had explicitly consented to the use of their photographs for research purposes.

⁵Predicting body weight from a facial image has been used in earlier experiments (Logg et al., 2019).

⁶This motive would not inherently bias the treatment-effect comparison, since it should be balanced across the *Punishment* and *No-Punishment* conditions, but it might otherwise dominate the delegation decision and may obscure the responsibility-shifting motive of interest.

each subject’s sense of her ability relative to the algorithm. Matching the algorithm to each subject neutralizes this motive, and making the match common knowledge prevents Player B from reading delegation as a signal of Player A’s believed relative competence. This closes off an important signal-extraction channel through which delegation might otherwise insulate Player A from punishment.

In the *Punishment* condition, Player A was further informed that Player B could reduce her payoff by up to £2, with the deduction depending on both her delegation decision and the realized outcome. In the *No-Punishment* condition, Player B had no such punishment option. The two delegation options (own decision / algorithm) were presented in random order across subjects. We attached no monetary cost to delegation. This brings us closer to Player A’s intrinsic preference over delegation in the *No-Punishment* condition, since with any positive cost, the responsibility-shifting motive would need to outweigh a price effect. Thus, costless delegation is standard in the literature (Bartling and Fischbacher, 2012; Coffman, 2011; Feier et al., 2022). If Player A chose to make the prediction herself, she entered her final prediction on the next screen. If she delegated instead, this screen was skipped.

We then asked Player A to report her beliefs about Player B’s punishment in each of the four scenarios resulting from the combination of delegation decision (yes/no) and prediction outcome (high/low payoff for Player B). Beliefs were elicited on the same scale as actual punishment (£0–£2 in £0.10 increments) and the four scenarios were presented in block-wise random order. We chose not to incentivize belief elicitation for three reasons. First, only two of the four scenarios remain payoff-relevant after the delegation decision, making the others purely hypothetical. Second, an incentivized mechanism could induce hedging behavior. Third, an incentive-compatible scoring rule across four scenarios would have substantially increased task complexity. Because we are nonetheless interested in comparing punishment beliefs to actual punishment, we included an attention check on the belief-elicitation screen. In the *No-Punishment* condition, no punishment was possible. Yet, we elicited beliefs about hypothetical punishment to maintain consistency across treatments.

Player A was then asked to assess her own performance in the first ten rounds by assigning a probability mass (summing to 100%) to each possible score from 0 to 10. We elicited this belief after the delegation decision, and after Player A had been told that the algorithm would perform equally well. While learning about algorithm performance may lead her to revise her beliefs about her own performance, it is precisely this post-update belief that is the analytically relevant input for any subsequent behavior. The elicitation was incentivized: Player A earned an additional £0.30 with the probability she had assigned to her actual score.⁷ Player A was also asked to estimate the distribution of scores in the population, assigning to each score from 0 to 10 the share of participants she expected to have achieved it. Combined with her belief about her own score, this locates where she places herself in the distribution and yields a measure of

⁷While this linear scoring rule is not proper—its payoff-maximizing report places all mass on the single most likely score, irrespective of risk attitude—in practice most subjects spread probability mass across several scores.

perceived relative ability. One score was drawn at random at the end of the experiment, and Player A earned £0.30 if her estimate for it fell within five percentage points of the realized population share.

Finally, we elicited risk preferences using an incentivized multiple price list in the style of [Holt and Laury \(2002\)](#) and administered a short questionnaire on socio-demographic characteristics and technology affinity, before presenting Player A with her preliminary payoff. Figure 12 in Appendix G summarizes the experimental timeline for Player A.

Player B

After providing informed consent, Player B was told that he had been randomly matched with a Player A and was introduced to the task Player A had completed. He learned that his own payoff would be determined by the outcome of Player A’s final prediction—£5 if it fell within 10 lbs of the true weight and £1 otherwise—while Player A’s payoff depended solely on her performance in the initial ten rounds. Crucially, Player B was also told that Player A had the option to delegate the prediction to an algorithm that performs equally well as Player A had performed in the earlier rounds, and that Player A was aware of this, making it common knowledge.

In the *Punishment* condition, Player B could choose to reduce Player A’s payoff by up to £2 in £0.10 increments. We focus on punishment rather than reward because the responsibility-attribution literature has predominantly operationalized blame through punishment ([Bartling and Fischbacher, 2012](#); [Coffman, 2011](#); [Steffel et al., 2016](#)). Imposing strictly positive punishment required Player B to forgo £0.10 of his own earnings, ensuring that punishment was costly to the punisher.⁸ We elicited punishment decisions for all four combinations of delegation (yes/no) and outcome (high/low payoff for Player B), implementing only the decision matching the realized combination. This strategy method gives us the within-subject power to detect punishment differences with the achieved sample size. Because the punishment decision is the central outcome variable for Player B, we included an attention check on the punishment-elicitation screen. In the *No-Punishment* condition, where actual punishment was not possible, we elicited hypothetical punishment decisions on the same scale to keep the experimental flow comparable across conditions. Player B’s perceptions of task difficulty, his risk preferences, and a short questionnaire followed, before he was presented with his final payoff. Figure 13 in Appendix G summarizes the experimental timeline for Player B.

Once Player B’s decisions had been recorded, Player A’s final payoff was calculated and all participants were paid on the day of participation.

⁸The cost is fixed rather than proportional to the punishment level.

Hypotheses and identification

Our design tests two main hypotheses, both grounded in the literature on responsibility avoidance through delegation (Bartling and Fischbacher, 2012; Coffman, 2011; Oexl and Grossman, 2013; Hamman et al., 2010; Hill, 2015; Steffel et al., 2016) and on responsibility attributions to algorithmic decision-makers (Feier et al., 2022; Kirchkamp and Strobel, 2019; Gogoll and Uhl, 2018).

H1 (Delegation): The introduction of potential punishment increases the likelihood that Player A delegates the decision to the algorithm.

H1 is grounded in empirical evidence that responsibility avoidance shapes delegation decisions involving human intermediaries. Bartling and Fischbacher (2012) also show that subjects are more likely to delegate morally consequential decisions even to a die roll—an early form of delegation to a non-intentional agent—suggesting that responsibility avoidance persists when the agent lacks intentionality. Whether the same logic extends to *algorithmic* delegates is less explored. Feier et al. (2022) report that delegation rates are similar for human and algorithmic intermediaries, suggesting that the *kind* of intermediary may matter less than the mere availability of one—but their design does not directly compare delegation rates across punishment regimes, leaving open the question H1 addresses. H1 is identified by the between-subject comparison of Player A’s delegation rate across the *Punishment* and *No-Punishment* conditions, which exploits the punishment manipulation as the only between-subject difference and holds fixed every other feature of the decision environment.

H2 (Punishment): Conditional on a low-payoff outcome for Player B, Player A is punished less when she delegated the decision to the algorithm than when she made it herself.

H2 mirrors H1 on the punishment side. Bartling and Fischbacher (2012) find that delegation to a randomization device still allows decision-makers to avoid some punishment, although less so than delegation to another human. Feier et al. (2022) find that decision-makers are rewarded less when a negative outcome is directly attributable to their own action, compared to when it results from an algorithm to which they delegated. However, our study differs from theirs in various respects: A different task, role separation between Player A and Player B, and the feature that their design does not control for relative performance at the individual level. As a result, in their setting the delegation decision carries a strong signal about Player A’s beliefs concerning her own ability relative to the algorithm. Thus, a failure to delegate may then be punished as overconfidence rather than for assuming responsibility, confounding inference on responsibility-shifting through pure intermediation with an inference about ability. H2 is identified by the within-subject comparison across the four (delegation x outcome) cells of Player B’s strategy-method punishment, which isolates the conditional-on-outcome effect of delegation on punishment, holding fixed both the realized consequences for Player B and all between-subject heterogeneity in Player B’s punishment proclivity.

Two further design features deserve note. First, although the prediction task affects Player B’s payoff, the incentives of the two players are nominally aligned: Player A has no direct financial interest in Player B receiving the low payoff, and a successful prediction strictly weakens her own expected punishment.⁹ [FB: maybe player A thinks ”no effort, no payoff” so that player B is undeserving. Do I have to talk about how incentives are aligned?] This alignment mirrors canonical real-world settings for algorithmic delegation—medical diagnosis, financial forecasting, judicial decision-support—in which the decision-maker has no direct material interest in the affected party’s outcome but bears reputational or moral consequences for poor outcomes. By replicating this alignment, our design speaks directly to the policy-relevant cases. A standard intention-based punishment motive should therefore have less bite here than in dictator-game settings where the decision-maker can choose to be selfish. Second, although the underlying decision (a prediction under uncertainty) is not itself moral, the setup nonetheless qualifies as part of the *moral domain* as defined in the algorithm-aversion literature (Gogoll and Uhl, 2018): Player A’s choice imposes monetary consequences on a matched recipient, and a poor decision causes real monetary harm. The implications of these two design choices for the interpretation of our results are taken up in Section 4. [FB: review if we take them up actually.]

4 Results

Data collection took place between February and June 2023 through the British online platform Prolific, with participation restricted to UK residents. A total of 322 subjects were recruited (161 assigned to Player A, 161 to Player B).¹⁰ On average, Player A spent 12 minutes completing the experiment and earned £3.63, while Player B spent approximately 8 minutes and earned £2.40. Treatment randomization was effective, as demonstrated by the balance across observable characteristics (Appendix A.1, Table 4).

Player A – Delegation behavior

We begin by examining the delegation behavior of Player A. In the *Punishment* condition, where Player B can punish, 42.5% of Player As delegated the decision to the algorithm. In the *No-Punishment* condition, where punishment is impossible, 59.3% delegated. Figure 3 illustrates the difference. The introduction of potential punishment thus significantly *reduces* the likelihood of delegation to the algorithm (Pearson Chi-squared, $p = 0.033$; Fisher’s exact, $p = 0.041$), which reverses the direction posited by *Hypothesis 1*.¹¹ [Comment for supervi-

⁹In the *Punishment* condition, a better outcome for Player B should reduce expected punishment for Player A. In the *No-Punishment* condition there is no punishment to reduce. Either way, Player A’s incentives are weakly aligned with Player B’s.


¹⁰We preregistered the collection of data from 320 participants (80 subjects per role per treatment). Due to attrition and subsequent re-matching we recruited an extra pair in the *No-Punishment* condition. Appendix D reports the full ex-ante power calculations underlying this sample size and distribution.

¹¹A within-subject corroborator points the same direction. Player As in the *No-Punishment* condition were asked, after their actual delegation decision, the hypothetical question of whether they would have delegated had

sor: Those are the p-values for the two-sided test. In the design section, we have one-sided hypotheses? Should we instead have equal delegation as the hypothesis and then reject that, particularly because the effect goes in the opposite direction as hypothesized, and because in the pre-registration the hypothesis was non-directional, even though there a direction was insinuated by the other hypothesis? In general, is it best practice to make the hypothesis rejectable, or do you not want to reject it and just put the hypothesis in a way that it actually reflect hypothesized behavior?] Delegation shares around 50% also suggest that subjects, on average, find the claim of equal relative performance between Player A and the algorithm conveyed by the design credible and their behavior reflects neither systematic preference for their own predictions nor systematic deference to the algorithm.

Result 1. *The possibility of punishment significantly reduces the likelihood that Player A delegates the decision to the algorithm.*

Figure 3: Delegation shares by condition



Note: Bars show the share of Player As who delegated the prediction for Player B to the algorithm in each condition. The two-sided difference is significant at the 5% level under both a Pearson Chi-squared test ($p = 0.033$) and a Fisher's exact test ($p = 0.041$).

Next, we examine the robustness of the observed treatment difference in the delegation decision. Specifically, we assess whether the difference in delegation rates remains when controlling for socio-demographic characteristics and Player A's actual and perceived task performance.

punishment been possible. Among the 81 respondents, 12 would have stopped delegating under hypothetical punishment (25% of delegators) while only 5 would have started (15.2% of non-delegators), constituting weakly significant switching behavior towards making the decision oneself with punishment possible (within-subject McNemar's test $p = 0.090$), and yielding a hypothetical delegation rate of 50.6%.

Table 1 reports estimates from a Probit model of the delegation decision. Column (1) reports the unconditional treatment effect. Column (2) adds a vector of socio-demographic and other background characteristics. Columns (3) and (4) add Player A’s actual performance (number of correct predictions in the first ten rounds) and her weighted-average belief about own performance, respectively. Column (5) restricts the sample to subjects who passed the attention check.¹² The treatment effect is robust across all five specifications: the coefficient on the No-Punishment indicator runs from 0.42 ($p = 0.034$) in the unconditional specification to 0.52 ($p = 0.025$) in the attention-check sample.¹³ Translating these into average marginal effects: removing the punishment possibility raises the probability of delegation by between 16.5 and 19.2 percentage points across all specifications. Neither actual nor perceived performance is significantly associated with delegation in the pooled sample, although actual performance loads negatively and approaches significance under the attention restriction ($p = 0.068$). We return to this pattern in Section 4.1.

Table 1: Determinants of the delegation decision

	Dependent variable: <i>Delegation</i>				<i>Passers</i>
	<i>Full sample</i>				
	(1)	(2)	(3)	(4)	(5)
No-Punishment indicator	0.423** (0.199)	0.435** (0.204)	0.472** (0.206)	0.448** (0.205)	0.517** (0.230)
Task performance			−0.119 (0.079)		−0.158* (0.086)
Perceived performance				−0.072 (0.056)	
Age		−0.002 (0.009)	−0.002 (0.009)	−0.002 (0.009)	0.002 (0.010)
Female		0.175 (0.227)	0.222 (0.228)	0.212 (0.229)	0.407 (0.254)
Socio-economic status		−0.038 (0.071)	−0.035 (0.072)	−0.031 (0.071)	−0.082 (0.079)
Went to uni		0.036 (0.227)	0.059 (0.225)	0.025 (0.227)	0.124 (0.247)
Technology score		−0.069 (0.092)	−0.048 (0.095)	−0.062 (0.092)	0.014 (0.107)
Leadership position		0.404* (0.225)	0.422* (0.226)	0.399* (0.225)	0.253 (0.253)
Constant	−0.189 (0.141)	0.018 (0.537)	0.242 (0.570)	0.268 (0.580)	0.134 (0.645)
AME of No-Punishment (pp)	16.5** (<i>p</i> = 0.026)	16.6** (<i>p</i> = 0.026)	17.8** (<i>p</i> = 0.016)	17.0** (<i>p</i> = 0.022)	19.2** (<i>p</i> = 0.017)
N	161	159	159	159	130
Pseudo <i>R</i> ²	0.020	0.039	0.049	0.047	0.064

Note: Coefficients from a Probit regression of the delegation decision on the No-Punishment indicator and the controls listed in each column. Robust (HC1) standard errors in parentheses. Significance: * $p < 0.10$; ** $p < 0.05$; *** $p < 0.01$ (two-sided). Two subjects are dropped in Columns (2)–(4) due to missing socio-demographic data. The *AME of No-Punishment* row reports the average marginal effect of the No-Punishment indicator on the probability of delegation, in percentage points.

¹²While the attention check was included on another screen, it may still correlate with who took the delegation decision seriously.


¹³Including a control for the random order in which the two delegation options appeared on screen leaves the treatment coefficient essentially unchanged (treatment coefficient 0.48 ($p = 0.019$), with the “algorithm first” order indicator at 0.27 ($p = 0.21$), starting from the specification in Column (3)).

Player B – Punishment behavior

We now turn to Player B’s punishment behavior, beginning with the *Punishment* condition, in which Player B was given a real, payoff-relevant punishment choice.

Punishment is substantial across all four (delegation, outcome) scenarios of the strategy method (Figure 4).¹⁴ Despite the sizable punishment after *good* outcomes, Player Bs punish bad outcomes more than good outcomes, by a mean difference of £0.21 (significant at the 1% level), indicating that measured punishment carries real informational content rather than being random. Substantial punishment even after a *good* outcome need not reflect noise either, since Player Bs can form views about the appropriateness of delegation, independently of the realized outcome, and can punish accordingly.¹⁵

Figure 4: Player B punishment by delegation and outcome scenario



Note: For each of the four (delegation, outcome) cells of the strategy method, the figure reports the average punishment chosen by Player B (left *y*-axis, in pounds; error bars are standard errors of the mean) and the share of Player Bs imposing non-zero punishment (right *y*-axis, in percent), in the *Punishment* condition (full sample $n = 80$, passers $n = 67$). The full sample is shown at full opacity. The attention-check passers are overlaid at lower opacity. Hypothetical punishment in the *No-Punishment* condition is reported in Appendix Figure 10.

Holding the outcome fixed, Player A is punished no differently for delegating to the algorithm than for deciding herself. After a good outcome, average punishment is £0.329 for delegated versus £0.297 for self-made decisions. After a bad outcome, £0.482 versus £0.553. Thus, the punishment gap for bad outcomes runs in the predicted direction but is small and statistically insignificant (£0.071; paired *t*-test $p = 0.13$, Wilcoxon signed-rank $p = 0.54$). The regression in Table 2 tells the same story: Punishment loads overwhelmingly on the outcome (bad-outcome coefficient 0.256, $p < 0.01$), the effect of delegation for good outcomes is essentially zero, and the *Delegated* $\times$ *Bad outcome* interaction—the differential that would capture any insulation from punishment for bad outcomes—is -0.103 and borderline significant ($p = 0.08$).¹⁶ The

¹⁴One-sample *t*-test: $p < 10^{-5}$ in every scenario.

¹⁵For example, Coffman (2011) likewise documents non-trivial punishment in settings where the decision-maker’s action could not have been improved upon.

¹⁶In the regression, because punishment for the different scenarios is elicited within-subject, any stable char-

conclusion is unchanged when we restrict to subjects who passed the attention check, which was included on the punishment elicitation screen (Table 2, column 2).

Table 2: Determinants of Player B’s chosen punishment

	Dependent variable: <i>Punishment</i> (£)	
	<i>Full sample</i>	<i>Passers</i>
	(1)	(2)
Delegated	0.032 (0.052)	0.001 (0.054)
Bad outcome	0.256*** (0.071)	0.260*** (0.080)
Delegated $\times$ Bad outcome	-0.103* (0.058)	-0.078 (0.052)
Player B belief about Player A perf.	-0.020 (0.032)	0.001 (0.037)
Age	0.002 (0.005)	-0.001 (0.006)
Female	-0.100 (0.128)	-0.064 (0.140)
Socio-economic status	0.083* (0.047)	0.072 (0.049)
Went to uni	-0.196 (0.140)	-0.098 (0.142)
Technology score	0.073 (0.053)	0.030 (0.059)
Leadership position	0.083 (0.142)	0.061 (0.164)
Constant	-0.077 (0.328)	0.020 (0.374)
Observations	320	268
Subjects	80	67
Adj. R^2	0.077	0.030

Note: OLS regression of Player B’s chosen punishment level on a delegation indicator, a bad-outcome indicator, their interaction, Player B’s belief about Player A’s performance, and a vector of socio-demographic controls. The sample is stacked over the four (delegation, outcome) cells of the strategy method (4 observations per Player B). The *Delegated* $\times$ *Bad outcome* interaction is the differential effect of delegation on punishment in the bad-outcome cell. Column (1) uses the full Punishment-condition sample; Column (2) restricts to Player Bs who passed the attention check on the punishment-elicitation screen, as preregistered. Standard errors clustered at the Player B level. Significance: * $p < 0.10$; ** $p < 0.05$; *** $p < 0.01$ (two-sided).

[Comment: Perhaps add something about the determinants of punishment in each cell or punishment difference for bad outcomes between delegated and self-made decisions? Old paragraph: ”Even though the algorithm and Player A are calibrated to perform identically, Player Bs may still hold heterogeneous beliefs about Player A’s *absolute* performance. We rule out the possibility that those beliefs drive cell-level punishment. Across all four scenarios, correlation between Player B’s weighted average belief about Player A’s score and chosen punishment are essentially zero ($|r| \leq 0.03$, all $p > 0.8$).”]

The same pattern holds on the extensive margin (Appendix Table 8). The share of Player Bs imposing *any* punishment is indistinguishable between delegated and self-made decisions for

acteristic of Player B is differenced out and cannot, by construction, confound the within-subject comparison. Thus, the control variables merely capture the between-subject association with Player B’s *average* punishment level. We find no striking relationship between Player B’s average punishment level and either their socio-demographic characteristics or their beliefs about Player A’s *absolute* performance.

each outcome (Figure 4, right axis). Punishment is also far from universal with 45.0% of Player Bs punishing in none of the four scenarios. Thus, within each scenario averages reflect a punishing minority rather than light, diffuse sanctioning.¹⁷

Result 2. *Player A does not avoid punishment by delegating the decision to the algorithm.*

Thus, Player B neither punishes delegated decisions substantially less harshly, nor less frequently. Therefore, in our experimental setting Player A cannot avoid punishment by delegating to the algorithm, nor is she punished more when delegating the decision to the algorithm. Instead, punishment is determined almost entirely by the realized outcome. This result also bears on *Result 1* as any treatment difference in delegation frequencies driven by punishment avoidance would be inconsistent with Player B’s observed punishment behavior.

[Comment: Hypothetical punishment shows similar results, but for bad outcomes the tendency reverses, i.e. delegated decisions would be punished more.]

4.1 Discussion of results

Result 1 does not merely fail to support *Hypothesis 1*, but the effect is statistically significant in the opposite direction. The possibility of punishment discourages, rather than encourages, delegation to the algorithm. Read together with Result 2, this is puzzling. The prospect of punishment changes Player A’s delegation decision, even though we detect no punishment penalty for delegating. In the remainder of this section, we examine the data more closely and weigh candidate explanations that could reconcile the two main results.

Departures from canonical settings. *Hypothesis 1* was grounded in evidence from settings that differ from ours in ways that plausibly matter. In the modified dictator games of [Bartling and Fischbacher \(2012\)](#); [Coffman \(2011\)](#); [Hamman et al. \(2010\)](#); [Oexl and Grossman \(2013\)](#), the decision-maker chooses between a fair and an unfair allocation. Thus, her intent is more transparent, the chosen allocation is implemented with certainty, and the unfair option benefits her directly at the recipient’s expense. Punishment in these settings sanctions a deliberate, self-serving act, and delegation is attractive precisely because it blurs the link between the principal and that act. Our prediction task departs from these games in all three respects. The outcome is uncertain, so a bad outcome may reflect bad luck rather than bad intent, and neither option affects Player A’s own monetary payoff, so the incentives of the two players can be considered aligned, leaving no self-serving choice to police.¹⁸ With little intent worth sanctioning, any intention-based punishment motive loses its bite—and with it, the value to Player A of placing an intermediary between herself and the outcome. Nevertheless, we chose this task deliberately

¹⁷Among the 44 Player Bs who punish in at least one scenario, sanctions are mostly non-trivial, so the costly option is used as a substantive penalty rather than a token gesture.

¹⁸We consider the incentives aligned even though Player A’s payoff is only indirectly affected by her decision, as long as a better outcome for Player B weakly reduces Player A’s expected punishment in the *Punishment* condition.

since prediction under genuine uncertainty, by a decision-maker with no material stake in the affected party’s outcome, is the environment in which real-world algorithmic delegation typically operates.

[FB comment: What about the more recent literature on Responsibility Avoidance through delegation to algorithms?]

This reasoning, however, explains why the predicted effect might vanish, not why it should reverse. The reversal calls for a motive that weighs against delegation when punishment is possible.

Anticipated differential punishment. The most natural candidate is anticipation of punishment. Player A may expect a delegated decision to be punished more harshly than a self-made one, and delegate less in response. Such an expectation would be consistent with evidence that observers view the delegation of morally consequential decisions to machines relatively critically (e.g., Gogoll and Uhl, 2018; Bigman and Gray, 2018; Jauernig et al., 2022; Bonnefon et al., 2024). Alternatively, the same anticipated punishment pattern could arise from a simple shirking motive, in which Player A expects greater punishment for shirking the effort to make a decision for Player B but, in the absence of punishment, can shirk without monetary consequences.¹⁹ Both motives carry the same testable implication that Player A expects a punishment penalty attached to delegation *as such*, conditional on the outcome.

Punishment beliefs. After the delegation decision but before learning the outcome, we elicited Player A’s beliefs about Player B’s punishment in each of the four (*delegation, outcome*) scenarios.

We interpret the belief measure with caution for several reasons. First, the elicitation was unincentivized and took place after the delegation decision, so reports may reflect hedging across scenarios or post-hoc rationalization rather than the beliefs that actually drove the choice. Second, hypotheticality may introduce additional noise as only two of the four scenarios remained relevant after Player A’s delegation decision in the *Punishment* condition, and all four were purely hypothetical in the *No-Punishment* condition. Anticipated punishment can affect behavior only where punishment is possible. We therefore restrict the analysis of punishment beliefs to the *Punishment* condition.


Two patterns emerge (Figure 5). First, Player As, on average, substantially overestimate punishment in every scenario. Second, Player As expect essentially the same punishment for delegated and self-made decisions, conditional on the outcome, correctly anticipating the pattern in Player B’s actual punishment behavior (Result 2).²⁰ Thus, measured beliefs provide

¹⁹While we observe punishment to be outcome-based, rather than delegation-based, a shirking motive — which would predict additional punishment for delegation — could at least rationalize the positive punishment observed after good outcomes under delegation.

²⁰The expected punishment penalty for delegation is small and statistically indistinguishable from zero in both outcome states, and a within-subject test finds no asymmetry between the two ($p = 0.376$ among attention

no support for a punishment anticipation mechanism for Result 1.²¹

Figure 5: Realized vs. anticipated punishment, by scenario



Note: Average punishment in pounds across the four (*delegation*, *outcome*) scenarios, attention-check passers only. Green bars: Player A’s anticipated punishment in the *Punishment* condition ($n = 61$). Magenta bars: Player B’s realized punishment in the *Punishment* condition ($n = 67$). Blue bars: Player A’s hypothetical punishment beliefs in the *No-Punishment* condition ($n = 70$), shown for calibration only — these beliefs concern a counterfactual that cannot materialize and do not enter Player A’s delegation decision in this arm. Whiskers denote ± 1 standard error.

Even though the belief measure was not incentivized, and at the time of the belief elicitation, two of the four scenarios could not materialize anymore, punishment beliefs carry within-subject information. Table 3 reports a Logit regression of the delegation decision on Player A’s belief difference (*no delegation* – *delegation*) in expected punishment, restricted throughout to the *Punishment* condition.²² Columns (1)–(3) use the simple average of the two belief differences (one for each outcome) as the single regressor: Column (1) reports the unconditional specification, Column (2) adds task-performance and socio-demographic controls, and Column (3) restricts the sample to subjects who passed the attention check on the belief-elicitation screen. Column (4) replaces the simple average with a weighted average, using Player A’s elicited probability of producing the high payoff as the weight—the statistic that enters an expected-utility calculation. Column (5) enters the two outcome-conditional differences separately.²³

check passers); see Appendix A.7.

²¹The same holds for the fully hypothetical beliefs elicited in the *No-Punishment* condition, which closely track those in the *Punishment* condition (Figure 5, blue bars). Thus, the manipulation appears not to have changed what subjects expected punishment to look like.

²²The two underlying belief measures compare expected punishment of a self-made and a delegated decision for each of the two possible outcomes for Player B, so that higher values correspond to relatively higher expected punishment for self-made decisions.

²³The two differences are nearly orthogonal in this sample (Pearson correlation -0.10 among passers), so the coefficients in Column (5) are additive partial slopes rather than competitors for the same variance. Under expected-utility reasoning, however, only their probability-weighted sum should enter the decision rule. We therefore treat Column (5) merely as a descriptive decomposition.

Table 3: Delegation and punishment beliefs

	Dependent variable: <i>Delegation</i>				
	<i>Full Punishment</i>		<i>Punishment, passers</i>		
	(1)	(2)	(3)	(4)	(5)
Punishment beliefs: simple average ($\bar{\Delta}$)	0.553 (0.656)	0.629 (0.744)	2.708** (1.286)		
Punishment beliefs: weighted by Pr(good)				3.329** (1.532)	
Punishment beliefs: bad outcome (no del. – del.)					0.937 (0.751)
Punishment beliefs: good outcome (no del. – del.)					1.579** (0.741)
Task performance		–0.337* (0.181)	–0.576** (0.231)	–0.544** (0.245)	–0.609*** (0.236)
Age		0.037* (0.020)	0.063** (0.029)	0.066** (0.030)	0.064** (0.030)
Female		0.213 (0.592)	1.522** (0.754)	1.630** (0.699)	1.334* (0.793)
Socio-economic status		–0.079 (0.185)	–0.012 (0.226)	–0.095 (0.233)	–0.003 (0.219)
Went to uni		0.108 (0.609)	–0.446 (0.712)	–0.257 (0.719)	–0.433 (0.707)
Technology score		–0.104 (0.230)	0.085 (0.290)	0.129 (0.295)	0.059 (0.283)
Leadership position		0.478 (0.510)	–0.485 (0.720)	–0.477 (0.731)	–0.387 (0.729)
Constant	–0.316 (0.228)	–0.536 (1.333)	–2.023 (2.080)	–2.153 (2.052)	–1.842 (2.103)
AME of belief measure (pp per £0.10)	1.3 ($p = 0.391$)	1.3 ($p = 0.395$)	4.6** ($p = 0.025$)	5.5** ($p = 0.018$)	
AME: good-outcome difference (pp per £0.10)					2.7** ($p = 0.018$)
AME: bad-outcome difference (pp per £0.10)					1.6 ($p = 0.215$)
Controls	×	✓	✓	✓	✓
Passers only	×	×	✓	✓	✓
N	80	78	60	60	60
Pseudo R^2	0.006	0.097	0.241	0.258	0.246

Note: Logit estimates with robust (HC1) standard errors in parentheses. Significance: * $p < 0.10$; ** $p < 0.05$; *** $p < 0.01$ (two-sided). Sample restricted to the *Punishment* condition throughout, since beliefs about a counterfactual that cannot materialise do not enter Player A’s decision rule in the *No-Punishment* arm. Belief differences are within-subject (*no delegation*) – (*delegation*) in expected punishment, so positive coefficients indicate that subjects who expect to be punished more harshly for not delegating (relative to delegating) are more likely to delegate. Column (1) uses the simple average of the two outcome-conditional belief differences without controls; Column (2) adds task performance and socio-demographic controls; Column (3) restricts to subjects who passed the attention check on the belief-elicitation screen, as preregistered; Column (4) replaces the simple average with a weighted average using Player A’s elicited probability of producing the high payoff ($\text{Pr}(\text{good}) = \text{wa.confidence}/10$); Column (5) replaces the single belief regressor with the two outcome-conditional belief differences entered separately. The *AME* rows report the average marginal effect of the belief measure entered in the respective column on the probability of delegation, in percentage points per £0.10 increase in the corresponding belief difference.

The estimated relationship is weak in the full sample but sharpens when restricting the sample to subjects who have passed the attention check on the belief elicitation screen. In Columns (1)–(2), expected differential punishment does not predict delegation. The coefficient is positive but not significant ($p > 0.39$), with an average marginal effect of 1.3 percentage points per £0.10 of belief difference. Among attention-check passers, the effect more than triples and becomes significant. A £0.10 larger expected punishment penalty for deciding herself is associated with a 4.6 percentage-point higher probability that Player A delegates (Column (3), coefficient 2.71, $p = 0.035$), and with 5.5 percentage points under the weighted specification (Column (4), 3.33, $p = 0.030$). In the decomposition, the good-outcome belief difference carries the relationship (1.58, $p = 0.033$, or 2.7 percentage points per £0.10), while the bad-outcome difference is positive but noisier (0.94, $p = 0.21$).

Taken together, the belief evidence is mixed. The measure itself is not pure noise. Within the *Punishment* condition, subjects whose beliefs embody a larger penalty for deciding themselves are more likely to delegate, so the delegation choice responds to belief asymmetries where they exist. But a margin whose average is zero provides no aggregate push. For anticipation of punishment to generate a 17 percentage-point delegation gap between treatments, Player As would need to expect systematically harsher punishment for delegated decisions. They do not and neither would such an expectation have been correct, as Player B’s realized punishment exhibits no delegation penalty either (Result 2). Anticipated differential punishment is thus a dimension along which subjects differ, but it cannot explain the treatment difference in delegation.


Selection into conditions. Nor is Result 1 an artifact of sample composition. Selection on observable characteristics cannot explain Result 1, as treatment randomization produced balance across socio-demographic and other background characteristics (Appendix A.1, Table 4). The same holds for performance. Average performance in the first ten rounds is balanced across conditions (2.76 vs. 3.09 correct predictions out of ten in the *Punishment* and *No-Punishment* conditions, respectively; two-sided t -test $p = 0.12$; Figure 6),²⁴ Player As in both conditions exhibit similar overconfidence about their own performance (4.65 vs. 4.87 in the *Punishment* and *No-Punishment* conditions, respectively; two-sided t -test $p = 0.47$),²⁵ and the treatment effect is robust to controlling for actual and perceived performance (Table 1, Columns (3)–(4)).

Performance and delegation. What differs across conditions is *who* delegates. A first indication already appears in the pooled delegation regression in Table 1, where performance loads negatively and approaches significance. This tendency crystallizes within the *Punishment* condition as the correlation between Player A’s score in the first ten rounds and the delegation

²⁴To the extent this small raw difference matters at all, it works against Result 1, because the *Punishment* condition performs slightly worse, and as we will see, it is worse performers who delegate more in that condition, so any residual imbalance would attenuate rather than generate the treatment difference.

²⁵Confidence is measured by the weighted-average belief about own performance. This overconfidence may partly reflect the information that the algorithm performs equally well.

Figure 6: Performance in the first ten rounds, by treatment



Note: Distribution of correct predictions (within 10 lbs of the true weight) over the first ten task rounds, by condition. Average performance is 2.93 correct out of ten and is statistically indistinguishable across conditions.


decision is -0.19 , with $p = 0.092$ (Figure 7),²⁶ while a comparable relationship is weaker in the *No-Punishment* condition.²⁷ Framed at the level of treatment differences, the delegation gap widens with performance. The composition of delegators shifts accordingly. Player As who delegated in the *Punishment* condition scored on average 2.47 out of ten in the first ten rounds, whereas delegators in the *No-Punishment* condition scored 3.06 (two-sided $p = 0.039$). The additional delegation in the *No-Punishment* condition thus comes disproportionately from the middle and upper part of the performance distribution. It is precisely these subjects who, when punishment is on the table, keep the decision for themselves.

This relationship is interesting because its natural explanation, namely that worse performers delegate because they believe the algorithm outperforms them, is ruled out in our design. In [Feier et al. \(2022\)](#), the algorithm's performance distribution is matched to the human distribution at the *population* level so that a high-performing subject who recognizes her position in that distribution rationally believes she outperforms the algorithm in expectation, and a low-performing one rationally believes the reverse. Delegation that declines with performance is then exactly what maximizing the expected outcome predicts. In our design, performance is equalized at the *individual* level and this is common knowledge. The pattern in the *Punishment*

²⁶Among attention-check passers the relationship is even stronger ($r = -0.22$, $p = 0.083$).


²⁷A similar relationship holds for perceived performance in the *Punishment* condition. Player A's weighted-average belief about her own performance enters with a small negative coefficient but is never statistically significant. Thus, actual performance discriminates sharply between delegators and non-delegators (2.38 vs. 3.00 correct out of ten, among attention-check passers) while perceived performance does not (4.69 vs. 4.71; Figure 8). Delegators under punishment are not subjects who *believe* they are worse than their algorithm, but they are subjects who *are* worse, and whose confidence does not reflect it.

Figure 7: Delegation share by performance and condition



Note: Each point reports the share of Player As at the corresponding first-ten-rounds score who delegated the final prediction. Dot size is proportional to the number of subjects at that score. The downward slope in the *Punishment* condition is the Pearson $r = -0.19$. The relationship is weaker in *No-Punishment*.

Figure 8: Performance and confidence distributions by treatment and delegation choice



Note: Restricted to attention-check passers. Within each treatment, the top panel shows the share of Player As at each first-ten-rounds score, separately for those who delegated (dark blue) and those who did not (light blue), with the weighted-average performance belief overlaid as a KDE (delegators solid magenta, non-delegators dashed pink). Boxplot strips below the top panel report the same four series. In the *Punishment* condition, actual performance discriminates between delegators and non-delegators (means 2.38 vs. 3.00) while perceived performance does not (means 4.69 vs. 4.71). In the *No-Punishment* condition, neither actual nor perceived performance discriminates (actual: 2.98 vs. 3.13; perceived: 4.85 vs. 5.33). The additional delegators in *No-Punishment* therefore come disproportionately from the middle and upper portion of the actual performance distribution.

condition must therefore reflect something other than beliefs about relative ability.

Interpreting the performance–delegation relationship. One reading of the relationship between performance and delegation could be a preference over *process*. Player As who have a good chance of producing the high payoff may prefer to bring it about themselves rather than through the algorithm — we refer to this as *process ownership* of the outcome. Equivalently, Player As with little chance of success may delegate to avoid being the proximate cause of a likely failure. Yet, while this motive can rationalize a negative relationship between performance and delegation, it would predict that relationship in *both* conditions and predict no treatment difference in delegation rates, unless the preference is *activated* by accountability, becoming behaviorally consequential once Player B’s stake in the decision is enforceable rather than merely material. Alternatively, Player A’s preferences for delegation could vary with the amount of *uncertainty* inherent in a decision. Dietvorst and Bharti (2020) document that the preference for algorithms declines as the uncertainty inherent in a decision rises. In our task, where most subjects score below five out of ten, better performance moves the success probability toward one half — arguably the point of maximal outcome uncertainty — so a distaste for delegating uncertain decisions would likewise depress delegation among better performers. However, also this channel applies to both conditions equally.

Effort to outperform the algorithm. Lastly, the prospect of punishment could motivate Player A to exert additional effort in making the decision for Player B. Despite the algorithm being calibrated to Player A’s past performance, her final prediction is not yet made. Thus, a Player A may believe that she can try especially hard and raise the chance of a good outcome for Player B above the algorithm’s, if the decision is not delegated. To the extent that Player A expects milder punishment after a good outcome, the prospect of punishment would then push her toward deciding herself. This possibility yields the prediction that non-delegators should improve their performance, specifically in the *Punishment* condition. Improvement in both conditions would instead indicate a treatment-independent (for instance, altruistic) motive, which cannot explain Result 1.

The data do not lend support to this story. Delegators do not make a final prediction themselves, so effort on the final prediction is observable merely for non-delegators. Measuring improvement as the difference between a subject’s average absolute prediction error over the first ten rounds and her absolute error on the final prediction, non-delegators in the *No-Punishment* condition improve on their record, whereas non-delegators in the *Punishment* condition do not. Because all Player As faced the same images in the first 10 rounds and on the final prediction, this differential cannot reflect task-specific variation. Realized outcomes tell the same story with 45.5% of non-delegators in the *No-Punishment* condition earning the high payoff for Player B, compared with 21.7% in the *Punishment* condition, even though the two groups’ performances are statistically indistinguishable over the first ten rounds (3.12 vs. 2.98 correct predictions; two-sided $p = 0.66$). If anything, effort seems to rise where punishment is impossible, which is

the opposite of what a punishment-motivated effort story would predict.

Taking stock. The reversal in Result 1 is robust, but its mechanism remains unresolved. It is not a composition artifact: the conditions are balanced in observables, performance, and confidence, and the treatment effect survives the corresponding controls. Two channels could, in principle, generate a treatment difference, but are both inconsistent with the data. Anticipated differential punishment would require a penalty for delegating that is absent from Player A’s beliefs on average — consistent with Player B’s realized punishment (Result 2) — even though belief heterogeneity predicts delegation at the margin. Punishment-induced effort predicts improvement concentrated in the *Punishment* condition, whereas improvement concentrates where punishment is impossible. What the data document instead is that delegation is withheld under the possibility of punishment disproportionately by better-performing Player As. Preferences over process ownership or over outcome uncertainty are consistent with this pattern, but treatment-independent and could explain Result 1 only if activated by accountability itself — an assumption our design cannot test, because the possibility of punishment and the presence of accountability coincide by construction.

4.2 Other

Order effects. The delegation decision and the punishment / belief-elicitation screens presented their options in randomised (or block-wise randomised) order. Order effects are localised to the *Punishment* condition: presenting the algorithm option first roughly doubles delegation (52.5% vs 32.5% when the own-decision option is first; $\chi^2 p = 0.070$), while the *No-Punishment* condition shows no order effect ($p = 0.674$); the treatment $\times$ order interaction is not significant ($p = 0.115$), so the H1 finding is not driven by interaction with order.

Power.

Add an ex-post (or, preferably, the original ex-ante) power analysis here, framed as in [Feier et al. \(2022\)](#).

Reconciling our findings with [Feier et al. \(2022\)](#). Two of our results sit in apparent tension with the closest precedent. First, [Feier et al. \(2022\)](#) document that decision-makers are punished *less* for algorithmically delegated decisions; we find no significant punishment differential. Second, the natural prediction from [Feier et al. \(2022\)](#)’s own findings is that the prospect of punishment should make algorithmic delegation *more* attractive; we find the reverse. We do not view these contrasts as contradictions, but as evidence that the design features which differ between the two studies are doing real work. The differential punishment in [Feier et al. \(2022\)](#) is most plausibly read as third parties extracting a signal from the delegation decision about the principal’s self-perceived ability relative to the algorithm—a signal that is

maximally informative when (as in their design) the algorithm is calibrated at the session level and when the principal’s relative ability is uncertain. Our individual-level calibration, made common knowledge, closes off this signal-extraction channel by construction: a Player B in our setting has no basis to read delegation as overconfidence, laziness, or ability avoidance. The disappearance of the punishment differential in our data is therefore consistent with, rather than refutes, the signal-extraction reading of [Feier et al. \(2022\)](#)’s result. By the same logic, the responsibility-shifting motive that [Feier et al. \(2022\)](#)’s setup leaves on the table is muted in ours, isolating a different and milder margin: the moral salience of accountability for the principal herself, independent of any third-party inference. The reversal we observe in the delegation decision speaks to that margin—and we are explicit that the question of which design feature is causally responsible for the sign reversal cannot be resolved without an experiment that varies relative-performance calibration directly, which we do not do here. [Comment: The second point is not a contradiction, because in Feier the two treatments that are compared are different to the two treatments that are compared in our paper.] [Comment: Issue here is that the relative performance concern applies to both treatments.]

We see scope, then, for the two studies to be parts of a single picture: [Feier et al. \(2022\)](#) identifies the signal-extraction channel through which delegation can shift third-party blame; we identify a process-ownership channel through which the prospect of accountability can pull the principal toward direct engagement with the decision. The two channels are not mutually exclusive. [Chevrier and Teixeira \(2024\)](#), looking at a third margin (the program behind the algorithm), and [Köbis et al. \(2025\)](#), looking at a fourth (the compliance asymmetry between human and machine agents on instructed unethical tasks), populate further nodes in this expanding map. A unified design that varies the relative-performance calibration, the punishment possibility, and the visibility of the algorithm’s compliance properties would be the natural next step. We discuss the implications for future work in Section 5.

[To resolve:]

- For two subjects, we must be missing realizations of control variables or something, because the sample size decreases from 161 to 159. Do we leave this unaddressed or should we include a footnote?
- Inequality/deservingness motive: Player A may think that Player B does not have to work to potentially earn 5 pounds and consider this unfair. However, Player A does not know what else Player B did have to do. Player B may want to bring down Player A’s payoff so that they earn similarly, but also here there is uncertainty and this would just be outcome-based and not prescribe a treatment effect [Comment: Do not discuss this.]
- Is there any benefit in looking at hypothetical punishment of the No-Punishment treatment? Given that there are a lot of hypotheticals also in the Punishment treatment, should we perhaps not even pool for any analysis that only looks at anticipated punishment, however not when relating the delegation decision to anticipated punishment

[Comment: Deacuse if it is entirely hypothetical, why would the beliefs matter for the delegation decision?]; How can we exploit that at the belief elicitation stage, a subset of scenarios are already for sure not relevant, e.g. we could check whether the difference between good and bad for delegated decisions depends on whether the decision was delegated, and the same for non-delegated decisions.

- Similarly to Player A, Player B overestimates Player A's task performance. Therefore, the likelihood with which the good outcome scenario materializes is anticipated to be higher than it actually is. Can we do anything with this? [Comment: If Player B thinks the task is easy, they could punish more for bad outcomes. Punishment (all four individually, but also differences) is independent (Pearson and Kendall) of beliefs about Player A's performance, so this holds consistently across the whole (mostly supported between 2 and 6) distribution of performance beliefs. Hence, from a punishment perspective it doesn't matter whether you delegate when task is easy or harder. For those thinking longer about the punishment decision, there is a positive correlation between easiness of the task and punishment for a bad outcome when you don't delegate vs. when you delegate.] Of course, the perception is again tainted by Player B being told that performance was the same.
- What can we do with the hypothetical delegation choice of subjects in the No-Punishment treatment? Empirically, the result moves into the same direction (ttest p-value: 0.089)]
- Are risk attitudes related to any decisions? More risk-seeking leading to less delegation maybe or what could be the hypothesis? [Comment: Could be interpreted as a coefficient as part of the main result 1 regression.] More risk-seeking people could also risk being punished more, but I do not see how that would relate to the results. [Comment: Result in Feier et al.: risk averse subjects were also more likely to delegate.]
- Do robustness analysis with performance not in points, but in terms of distance from the truth.
- There are hedging motives in the performance elicitation task, but let's ignore those.
- Ways to filter: (i) failed attention checks, (ii) did not spend a lot of time on the page (delegation page and beliefs about punishment/punishment page)
- Player B: Is punishment correlated with any socio-demographics. Should we not do a regression here as well? Related to this: For those that state that individuals should be held responsible for AI-generated decisions responsibility avoidance is less possible (pos. correlation between 'nodel-del-bad' and 'responsibility') (maybe should recode variable 'responsibility').
- Main result 1 is consistent with Gogoll and Uhl 2018: Algorithm use was punished more.
- Check: In python, use `ttest_rel` instead of `ttest_ind` to compare punishment.

- Check: Feier et al. use GLM Logit IRLS. Is that different from what I am using?
- Maybe check if Player As that are better/more confident in last guess have different beliefs.
- Scrutinize Feier et al. analysis, e.g. They have the same roles: Is punishment behavior independent of delegation behavior? Is it in line with delegation behavior?
- When showing balance across observables, then be careful to use chi-squared test for categorical variables.
- Show that randomization also worked on time taken, etc?
- For each analysis, make explicit whether this is with the reduced sample (attention checks) or full sample. Always consider also presenting the robustness check with the other.

[Design issue: Why not have punishment beliefs before delegation decision?] [Design question: Would it have been interesting to compare the delegation behavior between a decision for oneself vs. someone else?]

[Consideration for literature review: Regarding Feier et al: If I always say that punishment in their experiment could be due to the signal value in delegation (concerns regarding relative performance), then this holds for both their treatments so should technically not explain why they find differential rewarding in the machine treatment.]

5 Discussion and Conclusion

This paper asks two questions about how the prospect of punishment interacts with the use of algorithms as decision aids. Can decision-makers shift responsibility for bad outcomes onto algorithms by delegating their decisions? And does the anticipation of being held responsible motivate delegation in the first place? We address both questions in a preregistered online experiment in which Player A makes a payoff-relevant prediction for Player B and may delegate the prediction to an algorithm calibrated to perform as well as Player A herself. A between-subject manipulation switches off punishment by Player B in the *No-Punishment* condition; the *Punishment* preserves the natural setting in which the prospect of punishment is on the table.

Two findings emerge. First, and most striking, the *prospect* of punishment *reduces*, rather than increases, the propensity to delegate to the algorithm: 42.5% of Player As delegate when punishment is possible, compared to 59.3% when it is not. This direction is the opposite of what the responsibility-avoidance literature would predict, and we view the reversal as the paper’s main empirical contribution—noting that, because the result reverses the preregistered direction, the analyses we use to interpret it (Section 4.1) are exploratory rather than confirmatory. Second, conditional on the outcome for Player B, we cannot reject the hypothesis that delegated and

self-made decisions are punished equally. The point estimate runs in the predicted direction (lower punishment under delegation in the bad-outcome cell) but is small relative to the test’s confidence interval at the achieved sample size; we therefore frame Result 2 as a failure to detect insulation rather than as positive evidence that delegation cannot insulate. Player A does not appear to be able to use this particular algorithm as a meaningful blame shield in the bounded sense the test can assess.

Mechanism. Why does punishment reduce, rather than increase, delegation? The data are inconsistent with a simple anticipation mechanism: Player As do not on average expect harsher punishment for delegated decisions, and they correctly anticipate the absence of differential punishment that we observe in Player B’s actual choices. The data are likewise inconsistent with the simple form of an effort story in which the prospect of punishment motivates Player A to outperform the algorithm: among non-delegators, those in the *No-Punishment* condition (where punishment is impossible) outperform those in the *Punishment* condition, the reverse of what such a story implies. The reading most consistent with the data is one of *process ownership*: Player As who actively choose not to delegate take responsibility for the outcome, exert correspondingly higher effort, and that share is higher when punishment is on the table. The supporting evidence is exploratory rather than confirmatory, however, and rests on a single performance-by-delegation heterogeneity pattern in a 60-subject sub-sample (Column (3) of Table 1); we cannot, with this design, separate process ownership cleanly from an experimenter-demand reading in which the punishment cue itself nudges Player A toward direct engagement. We therefore present process ownership as the most consistent reading of the data we have, not as a confirmed mechanism. Disentangling it from demand is a clean direction for follow-up work, sketched below.

This refinement is consistent with, but goes beyond, the standard reading of the responsibility-avoidance literature. The classic finding—that delegation reduces punishment of the principal and that this prospect motivates delegation—was established in dictator-game settings where the underlying choice is between a known fair and a known unfair allocation, the intermediary is human, and the principal can be held accountable specifically for choosing a particular intermediary (Bartling and Fischbacher, 2012; Coffman, 2011; Oexl and Grossman, 2013). When the intermediary is an algorithm calibrated to mimic the principal and the underlying decision involves genuine uncertainty, none of the standard mechanisms through which delegation can shift responsibility—signal value of the delegation choice, the human-human punishment channel, the certainty of the unfair outcome—are available. What remains is a baseline outcome-based punishment, supplemented by an asymmetric desire to be the proximate cause of good outcomes when the moral stakes are salient.

Implications, with explicit bounds. Our findings speak to the design of accountability regimes governing algorithmic decision-making. The European Union’s AI Act (European Union, 2024) and analogous regulatory frameworks elsewhere place increasing emphasis on as-

signing responsibility to human users of high-risk algorithmic systems. The implicit theory of behavior underlying these provisions is that holding humans accountable for algorithm-mediated decisions *disciplines* their use of those algorithms. Our results offer a behavioral complement to this view: when punishment is on the table, decision-makers in our setting do not simply delegate-and-blame the algorithm; they choose to take direct responsibility for the decision more often. From the perspective of efficiency, this need not be welfare-increasing: in domains where algorithms genuinely outperform humans, holding humans liable for outcomes may push decision-makers away from the more accurate technology. The trade-off between encouraging algorithmic adoption and assigning meaningful responsibility is, on this reading, sharper than the existing literature has suggested.

We note explicitly the bounds of this reading. Our evidence is a single online experiment with 322 UK Prolific participants on a stylised weight-prediction task, with stakes of £5 for Player B and a punishment range of £0–£2. The high-stakes algorithmic domains that motivate the regulatory debate—medical diagnosis, judicial risk assessment, hiring, credit allocation—differ from our setting in stake magnitude, professional context, expert training, the institutional structure of accountability, and the kind of population at the decision margin. We therefore view the result as identifying a behavioral channel that exists in principle and is detectable in our setting, not as a calibration of how big a lever accountability is in any specific applied domain. Mapping the magnitude of the effect against external benchmarks—e.g., physician adoption of diagnostic algorithms under malpractice exposure, or loan-officer deviation from credit-scoring algorithms under fair-lending audit—would require either domain-specific replication or instruments to which we do not have access. We flag this gap explicitly so that readers can calibrate the policy implication appropriately.

Limitations. Several caveats are in order. The belief measure we elicit is unincentivised and noisy; it conveys ordinal patterns reliably but cannot support fine-grained quantitative claims about Player A’s anticipations. Two of the four belief scenarios are hypothetical conditional on the delegation decision, and all four are hypothetical in the *No-Punishment* condition; this further limits the inference, and we plan to report the H2 patterns separately for the payoff-relevant and the hypothetical cells once the analysis pipeline is in place. We use the strategy method for Player B’s punishment decision, which is standard in this literature but may yield different absolute levels of punishment than a direct elicitation. Player A also makes the final prediction whether or not she ultimately decides to delegate, which equates the time cost of the two options but may signal her commitment differently than a design in which the prediction is skipped under delegation. We did not include a separate manipulation-comprehension check verifying that Player A internalised the punishment-on/off distinction; we relied on the prominent screen-level framing and on attention-check pass rates as indirect evidence, but a future replication should add this check. The absolute stakes are modest (£5 high payoff, £0–2 punishment range): we cannot speak to how the effect scales with stakes, and a stake-magnitude robustness check is the first item in the future-research list below. Our sample is

drawn from UK Prolific participants, which limits external validity along the dimensions in which Prolific samples differ from broader populations and from professional decision-makers in the high-stakes algorithmic domains we discuss above.

A natural concern is whether the punishment possibility cues Player A to behave more conservatively for reasons unrelated to the substantive question—an experimenter-demand effect. We cannot rule this concern out from the data. The direction of any such effect is also ambiguous: a decision-maker who is reminded of being judged could plausibly delegate *more* (to shed responsibility) or *less* (to demonstrate effort and commitment); the literature does not unambiguously favor one over the other. We note two patterns weakly inconsistent with a strong demand reading—behavior does not track elicited beliefs as tightly as a pure-demand story would predict, and the sharp performance heterogeneity in the *Punishment* condition (better performers delegating less) is harder to motivate from demand alone—but we treat these as suggestive rather than dispositive. A direct test in the spirit of [de Quidt et al. \(2018\)](#) would be the natural next step.

Directions for future research. Several extensions would strengthen the inference. A direct-method companion would test the robustness of Result 2 against the strategy-method elicitation, and a meaningfully larger sample would let us put a tight confidence interval on the small but persistent point estimate of insulation under delegation. A condition in which Player A is told explicitly that Player B is unaware of the equal-performance calibration would test whether the relative-performance signal is the binding margin distinguishing our results from [Feier et al. \(2022\)](#). Conditions that experimentally vary the algorithm’s performance relative to Player A (above, equal, below) would map the generality of the process-ownership reading. A demand-effect probe along the lines of [de Quidt et al. \(2018\)](#) would put a credible upper bound on the contribution of M4. A stake-magnitude robustness check, varying the £5/£1 high/low payoffs and the £0–£2 punishment range, would speak to whether the effect scales with stakes in the direction that policy-relevant settings would imply. Finally, replication in a different country, in a different decision domain (e.g., a medical-triage or hiring vignette), or with a sample of professional decision-makers rather than online crowd-workers would speak directly to the external-validity bounds we describe above.

Taken together, our results suggest a more textured story than the one that has emerged from the existing literature on responsibility avoidance and algorithm aversion. Within the bounds of the setting we study, the capacity to use algorithms as a tool for shifting responsibility appears limited, and the prospect of being held responsible—rather than encouraging delegation as a hedge—appears to push decision-makers towards direct engagement with the consequential choice. Whether the same channel operates in the high-stakes algorithmic domains that animate the regulatory debate is an empirical question that this study can frame but cannot, on its own, answer.

References

- Aghion, Philippe and Jean Tirole**, “Formal and Real Authority in Organizations,” *Journal of Political Economy*, 1997, 105 (1), 1–29.
- Bartling, B. and U. Fischbacher**, “Shifting the Blame: On Delegation and Responsibility,” *The Review of Economic Studies*, January 2012, 79 (1), 67–87.
- Bartling, Björn, Urs Fischbacher, and Simeon Schudy**, “Pivotality and responsibility attribution in sequential voting,” *Journal of Public Economics*, 2015, 128, 133–139.
- Berger, Benedikt, Martin Adam, Alexander Rühr, and Alexander Benlian**, “Watch Me Improve — Algorithm Aversion and Demonstrating the Ability to Learn,” *Business & Information Systems Engineering*, 2021, 63 (1), 55–68.
- Bigman, Yochanan E. and Kurt Gray**, “People Are Averse to Machines Making Moral Decisions,” *Cognition*, 2018, 181, 21–34.
- , **Desman Wilson, Mads N. Arnestad, Adam Waytz, and Kurt Gray**, “Algorithmic Discrimination Causes Less Moral Outrage Than Human Discrimination,” *Journal of Experimental Psychology: General*, 2023, 152 (1), 4–27. VERIFY middle author spellings against the APA page.
- Bockstedt, Jesse C. and Joseph R. Buckman**, “Humans’ Use of AI Assistance: The Effect of Loss Aversion on Willingness to Delegate Decisions,” *Management Science*, 2025, 72 (1), 323–342.
- Bogert, Eric, Aaron Schechter, and Richard T. Watson**, “Humans Rely More on Algorithms than Social Influence as a Task Becomes More Difficult,” *Scientific Reports*, 2021, 11, 8028.
- Bonnefon, Jean-François, Iyad Rahwan, and Azim Shariff**, “The Moral Psychology of Artificial Intelligence,” *Annual Review of Psychology*, 2024, 75, 653–675.
- Burton, Jason W., Mari-Klara Stein, and Tina Blegind Jensen**, “A systematic review of algorithm aversion in augmented decision making,” *Journal of Behavioral Decision Making*, April 2020, 33 (2), 220–239.
- Castelo, Noah, Maarten W. Bos, and Donald R. Lehmann**, “Task-Dependent Algorithm Aversion,” *Journal of Marketing Research*, 2019, 56 (5), 809–825.
- Chen, Daniel L., Martin Schonger, and Chris Wickens**, “oTree — An Open-Source Platform for Laboratory, Online, and Field Experiments,” *Journal of Behavioral and Experimental Finance*, 2016, 9, 88–97.

- Chevrier, Mathieu and Vincent Teixeira**, “Algorithm Delegation and Responsibility: Shifting Blame to the Programmer?,” 2024. GREDEG Working Paper No. 2024-04, Université Côte d’Azur.
- Chugunova, Marina and Daniela Sele**, “We and It: An interdisciplinary review of the experimental evidence on how humans interact with machines,” *Journal of Behavioral and Experimental Economics*, August 2022, *99*, 101897.
- Coffman, Lucas C**, “Intermediation Reduces Punishment (and Reward),” *American Economic Journal: Microeconomics*, November 2011, *3* (4), 77–106.
- Dawes, Robyn M.**, “The Robust Beauty of Improper Linear Models in Decision Making,” *American Psychologist*, 1979, *34* (7), 571–582.
- de Quidt, Jonathan, Johannes Haushofer, and Christopher Roth**, “Measuring and Bounding Experimenter Demand,” *American Economic Review*, 2018, *108* (11), 3266–3302.
- Dietvorst, Berkeley J. and Soham Bharti**, “People Reject Algorithms in Uncertain Decision Domains Because They Have Diminishing Sensitivity to Forecasting Error,” *Psychological Science*, 2020, *31* (10), 1302–1314.
- , **Joseph P. Simmons, and Cade Massey**, “Algorithm aversion: People erroneously avoid algorithms after seeing them err.,” *Journal of Experimental Psychology: General*, 2015, *144* (1), 114–126.
- , – , and – , “Overcoming Algorithm Aversion: People Will Use Imperfect Algorithms If They Can (Even Slightly) Modify Them,” *Management Science*, 2018, *64* (3), 1155–1170.
- Erat, Sanjiv**, “Avoiding Lying: The Case of Delegated Deception,” *Journal of Economic Behavior & Organization*, 2013, *93*, 273–278.
- European Union**, “Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 Laying Down Harmonised Rules on Artificial Intelligence (Artificial Intelligence Act),” Official Journal of the European Union 2024. VERIFY citation form.
- Feier, Till, Jan Gogoll, and Matthias Uhl**, “Hiding Behind Machines: Artificial Agents May Help to Evade Punishment,” *Science and Engineering Ethics*, April 2022, *28* (2), 19.
- Fershtman, Chaim and Uri Gneezy**, “Strategic Delegation: An Experiment,” *The RAND Journal of Economics*, 2001, *32* (2), 352–368.
- Fiorina, Morris P.**, “Legislator Uncertainty, Legislative Control, and the Delegation of Legislative Power,” *Journal of Law, Economics, & Organization*, 1986, *2* (1), 33–51.
- Freer, Mikhail, Daniel Friedman, and Simon Weidenholzer**, “Motives for Delegating Financial Decisions,” April 2024. arXiv:2309.03419 [econ.GN], v3 15 April 2024. University of Essex / UCSC.

- Gogoll, Jan and Matthias Uhl**, “Rage against the machine: Automation in the moral domain,” *Journal of Behavioral and Experimental Economics*, June 2018, *74*, 97–103.
- Hamman, John R, George Loewenstein, and Roberto A Weber**, “Self-Interest through Delegation: An Additional Rationale for the Principal-Agent Relationship,” *American Economic Review*, September 2010, *100* (4), 1826–1846.
- Hill, Adam**, “Does Delegation Undermine Accountability? Experimental Evidence on the Relationship Between Blame Shifting and Control,” *Journal of Empirical Legal Studies*, June 2015, *12* (2), 311–339.
- Holt, Charles A. and Susan K. Laury**, “Risk Aversion and Incentive Effects,” *American Economic Review*, 2002, *92* (5), 1644–1655.
- Hüholt, Nicola and Nora Szech**, “Trusting Machines with Morality – Delegating Moral Decisions to AI,” *European Economic Review*, 2026, *184*, 105255.
- Jauernig, Johanna, Matthias Uhl, and Gari Walkowitz**, “People Prefer Moral Discretion to Algorithms: Algorithm Aversion Beyond Intransparency,” *Philosophy & Technology*, 2022, *35* (1), 2.
- Kirchkamp, Oliver and Christina Strobel**, “Sharing responsibility with a machine,” *Journal of Behavioral and Experimental Economics*, June 2019, *80*, 25–33.
- Köbis, Nils, Zoe Rahwan, Raluca Rilla, Bramantyo Ibrahim Supriyatno, Clara Bersch, Tamer Ajaj, Jean-François Bonnefon, and Iyad Rahwan**, “Delegation to artificial intelligence can increase dishonest behaviour,” *Nature*, 2025, *646* (8083), 126–134.
- Litterscheidt, Roman and David J. Streich**, “Financial Education and Digital Asset Management: What’s in the Black Box?,” *Journal of Behavioral and Experimental Economics*, 2020, *87*, 101573.
- Logg, Jennifer M., Julia A. Minson, and Don A. Moore**, “Algorithm appreciation: People prefer algorithmic to human judgment,” *Organizational Behavior and Human Decision Processes*, March 2019, *151*, 90–103.
- Maasland, Christian and Kristina S. Weißmüller**, “Blame the Machine? Insights From an Experiment on Algorithm Aversion and Blame Avoidance in Computer-Aided Human Resource Management,” *Frontiers in Psychology*, May 2022, *13*, 779028.
- Mahmud, Hasan, A.K.M. Najmul Islam, Syed Ishtiaque Ahmed, and Kari Smolander**, “What influences algorithmic decision-making? A systematic literature review on algorithm aversion,” *Technological Forecasting and Social Change*, February 2022, *175*, 121390.
- Meehl, Paul E.**, *Clinical versus Statistical Prediction: A Theoretical Analysis and a Review of the Evidence*, Minneapolis: University of Minnesota Press, 1954.

- Naquin, Charles E. and Terri R. Kurtzberg**, “Human reactions to technological failure: How accidents rooted in technology vs. human error influence judgments of organizational accountability,” *Organizational Behavior and Human Decision Processes*, 2004, *93* (2), 129–141. VERIFY: original bib listed authors in reverse order (Kurtzberg, Naquin) and a different title. The above is the only 2004 Naquin-Kurtzberg paper on technology and attributions that the verifier could locate. CONFIRM with Felix.
- Oexl, Regine and Zachary J. Grossman**, “Shifting the blame to a powerless intermediary,” *Experimental Economics*, September 2013, *16* (3), 306–312.
- Pezzo, Mark V. and Stephanie P. Pezzo**, “Physician Evaluation After Medical Errors: Does Having a Computer Decision Aid Help or Hurt in Hindsight?,” *Medical Decision Making*, 2006, *26* (1), 48–56. VERIFY.
- Qin, Xin, Xiang Zhou, Chen Chen, Dongyuan Wu, Hansen Zhou, Xiaowei Dong, Limei Cao, and Jackson G. Lu**, “AI Aversion or Appreciation? A Capability–Personalization Framework and a Meta-Analytic Review,” *Psychological Bulletin*, 2025, *151* (5), 580–599.
- Schelling, Thomas C.**, *The Strategy of Conflict*, Cambridge, MA: Harvard University Press, 1960.
- Shank, Daniel B., Alyssa DeSanti, and Timothy Maninger**, “When are Artificial Intelligence versus human agents faulted for wrongdoing? Moral attributions after individual and joint decisions,” *Information, Communication & Society*, 2019, *22* (5), 648–663.
- Sharan, Navya Nishith and Daniela Maria Romano**, “The Effects of Personality and Locus of Control on Trust in Humans versus Artificial Intelligence,” *Heliyon*, 2020, *6* (8), e04572.
- Steffel, Mary, Elanor F. Williams, and Jaclyn Permann-Graham**, “Passing the buck: Delegating choices to others to avoid responsibility and blame,” *Organizational Behavior and Human Decision Processes*, July 2016, *135*, 32–44.
- Stein, Jens-Peter et al.**, “[NOT FOUND] Don’t Let Me Down: The Role of Algorithmic Complexity in Algorithm Aversion,” *Working Paper*, 2020. VERIFY: claim-verifier could not locate this paper in any indexed database. Title and authors may be partially fabricated by an earlier session. Original source unknown. RECOMMEND: drop the citation in literature.tex or replace with a verified algorithm-complexity-aversion reference (e.g., Castelo, Bos & Lehmann 2019 on task-dependent algorithm aversion).
- Sunstein, Cass R. and Jared H. Gaffe**, “An Anatomy of Algorithm Aversion,” *Columbia Science and Technology Law Review*, 2024, *26*, 290–319. VERIFY volume/pages against the published issue.

- Tobia, Kevin, Aileen Nielsen, and Alexander Stremitzer**, “When Does Physician Use of AI Increase Liability?,” *Journal of Nuclear Medicine*, 2021, *62* (1), 17–21.
- Tontrup, Stephan and Christopher Jon Sprigman**, “Strategic Delegation of Moral Decisions to AI,” October 2025. SSRN Working Paper No. 5696827, posted 6 October 2025.
- Weaver, R. Kent**, “The Politics of Blame Avoidance,” *Journal of Public Policy*, 1986, *6* (4), 371–398.

Appendix

A Additional Tables

A.1 Condition balance — Player A

Table 4: Condition balance — Player A

Variable	<i>Punishment</i>	<i>No-Punishment</i>	<i>p-value</i>
Age	39.026	37.877	0.552
Female	0.487	0.494	1.000
Socio-economic status	5.179	5.321	0.564
Went to uni	0.575	0.556	0.928
Technology score	2.087	2.284	0.322
Leadership position	0.362	0.333	0.824

Note: Mean values by condition for Player A, with p -values from two-sided t -tests for continuous variables and Chi-squared tests for binary variables. *Socio-economic status* is self-reported on a 1–10 scale; *Went to uni* indicates a university degree; *Technology score* is constructed from weekly device usage, programming skills, cryptocurrency knowledge, technology use at work, and similar items; *Leadership position* indicates self-reported management or supervisory experience.

Complementary to the variable-by-variable comparisons in Table 4, we test their joint significance by regressing the treatment indicator on all characteristics in a linear probability model. The observables are jointly uninformative about treatment ($F(6, 152) = 0.36$, $p = 0.90$; $R^2 = 0.01$; $N = 159$). A logit specification yields the same conclusion (likelihood-ratio $\chi^2(6) = 2.26$, $p = 0.89$).

We do not separately report a balance table for Player B. Player B’s punishment behavior is identified within-subject across the four scenarios of the strategy method, so balance across conditions is not strictly required for the punishment analysis.

A.2 Full regression results — delegation decision

A.3 Robustness — delegation decision

A.4 Determinants of Player B’s chosen punishment

A.5 Robustness — Player B’s chosen punishment

A.6 Extensive margin of punishment

Table 8 re-estimates the punishment regression on the *extensive* margin, replacing the punishment amount with an indicator for whether Player B imposed any punishment in a given cell, estimated by probit. The conclusion matches the intensive-margin result in the body. A bad

Table 5: Determinants of the delegation decision — full controls

	Dependent variable: <i>Delegation</i>				<i>Passers</i>
	<i>Full sample</i>				
	(1)	(2)	(3)	(4)	(5)
No-Punishment indicator	0.677** (0.320)	0.696** (0.329)	0.757** (0.334)	0.721** (0.333)	0.834** (0.376)
Task performance			−0.192 (0.128)		−0.254* (0.142)
Perceived performance				−0.118 (0.094)	
Age		−0.004 (0.015)	−0.004 (0.015)	−0.003 (0.015)	0.002 (0.017)
Female		0.279 (0.366)	0.350 (0.369)	0.332 (0.372)	0.651 (0.414)
Socio-economic status		−0.062 (0.113)	−0.058 (0.116)	−0.051 (0.115)	−0.134 (0.130)
Went to uni		0.061 (0.369)	0.091 (0.364)	0.047 (0.367)	0.184 (0.400)
Technology score		−0.110 (0.149)	−0.078 (0.153)	−0.101 (0.148)	0.021 (0.174)
Leadership position		0.647* (0.364)	0.676* (0.367)	0.641* (0.365)	0.413 (0.412)
Constant	−0.302 (0.226)	0.036 (0.862)	0.402 (0.930)	0.449 (0.945)	0.241 (1.062)
AME of No-Punishment (pp)	16.4** (<i>p</i> = 0.025)	16.5** (<i>p</i> = 0.026)	17.7** (<i>p</i> = 0.016)	16.9** (<i>p</i> = 0.022)	19.1** (<i>p</i> = 0.017)
N	161	159	159	159	130
Pseudo <i>R</i> ²	0.020	0.039	0.049	0.047	0.063

Note: Coefficients from a Logit regression of the delegation decision on the No-Punishment indicator and the controls listed in each column. Robust (HC1) standard errors in parentheses. Significance: * $p < 0.10$; ** $p < 0.05$; *** $p < 0.01$ (two-sided). Two subjects are dropped in Columns (2)–(4) due to missing socio-demographic data. The *AME of No-Punishment* row reports the average marginal effect of the No-Punishment indicator on the probability of delegation, in percentage points.

Table 6: Determinants of the delegation decision — robustness

	Dependent variable: <i>Delegation</i>		
	<i>Full sample</i>		<i>Punishment</i> (no excl.)
	(1)	(2)	(3)
No-Punishment indicator	0.771 (0.336)	0.677 (0.320)	
Task performance	−0.277 (0.149)		−0.395 (0.185)
Overconfidence	−0.103 (0.094)		
Punishment beliefs: bad outcome (no del. − del.)			−0.277 (0.499)
Punishment beliefs: good outcome (no del. − del.)			0.812 (0.529)
Controls	yes	no	yes
Attention-check excl.	yes (Player A)	n/a	no
Sample	Full	Full	Punishment
N	159	161	78
Pseudo R^2	0.055	0.020	0.114

Note: Robustness specifications for the delegation logit. Column (1) adds *overconfidence* (defined as the weighted-average belief about own performance minus actual performance) to the main full-sample specification. The treatment effect remains: *No-Punishment indicator* $p = 0.022$. *Overconfidence* is not significantly associated with delegation $p = 0.272$. Column (2) reproduces the unconditional treatment effect for completeness ($p = 0.034$; identical to the main Col (1)). Column (3) restricts to the *Punishment* condition and includes the within-subject belief differences but *does not* exclude subjects who failed the belief-screen attention check; the heterogeneity result is preserved: *Task performance* $p = 0.033$; *Punishment beliefs: good outcome* $p = 0.125$. Robust HC1 standard errors in parentheses. Following AEA editorial policy, significance stars are not reported in the table.

outcome sharply raises the probability of punishment—an average marginal effect of about 18 percentage points ($p < 0.001$)—whereas delegation leaves it essentially unchanged: the average marginal effect of delegation is -1.8 percentage points and statistically insignificant ($p = 0.46$; -2.9 percentage points among attention-check passers, $p = 0.29$). Player A is therefore no less *likely* to be punished when she delegates, just as she is punished no less *harshly*.

A.7 Punishment beliefs and task performance

In the main text we claim that Player A’s punishment beliefs do not vary systematically with her own task performance, and that Player As do not expect a relatively larger punishment penalty for delegation in good-outcome states than in bad-outcome states. This appendix documents the supporting analyses, restricted to the *Punishment* condition and attention-check passers ($n = 61$) for consistency with the main within-subject regression.

Difference-of-differences in the expected delegation penalty. Define the expected punishment penalty for delegation in each outcome as $\Delta^{\text{good}} \equiv \beta^{\text{del,high}} - \beta^{\text{self,high}}$ and analogously Δ^{bad} . A within-subject paired t -test of $H_0 : \Delta^{\text{good}} = \Delta^{\text{bad}}$ yields $\bar{\Delta}^{\text{good}} - \bar{\Delta}^{\text{bad}} = +0.075$ ($t = +0.99$, $p = 0.325$) on the full Punishment sample and $+0.085$ ($t = +0.89$, $p = 0.376$) on the passers. We cannot reject equality of the two delegation penalties.

Table 7: Determinants of Player B’s chosen punishment — robustness

	Dependent variable: <i>Punishment</i>	
	<i>Excl.</i>	<i>Incl. att. failers</i>
	(1)	(2)
Delegated	0.001 (0.054)	0.032 (0.052)
Bad outcome	0.260 (0.080)	0.256 (0.071)
Delegated $\times$ Bad outcome	−0.078 (0.052)	−0.103 (0.058)
Player B belief about Player A perf. (wa)	0.001 (0.037)	−0.020 (0.032)
Age	−0.001 (0.006)	0.002 (0.005)
Female	−0.064 (0.140)	−0.100 (0.128)
Socio-economic status	0.072 (0.049)	0.083 (0.047)
Went to uni	−0.098 (0.142)	−0.196 (0.140)
Technology score	0.030 (0.059)	0.073 (0.053)
Leadership position	0.061 (0.164)	0.083 (0.142)
Constant	0.020 (0.374)	−0.077 (0.328)
Attention-check excl.	yes	no
N (subject-by-cell)	268	320
Adj. R^2	0.030	0.077

Note: Punishment regression with Player B’s belief about Player A’s performance (*wa_difficulty*) added as a regressor. Column (1) excludes Player Bs who failed the punishment-screen attention check (preregistered); Column (2) keeps them. Standard errors clustered at the Player B level. The delegation, outcome, and interaction coefficients are stable across both specifications, and the belief regressor is not significantly associated with chosen punishment. H2 paired-test on the bad-outcome cell: with exclusion paired- t $p = 0.153$, Wilcoxon $p = 0.533$ ($n = 67$); without exclusion paired- t $p = 0.131$, Wilcoxon $p = 0.544$ ($n = 80$). Following AEA editorial policy, significance stars are not reported.

Table 8: Extensive margin of punishment — probit for $\Pr(\text{punishment} > 0)$


	Dependent variable: <i>Any punishment</i>	
	<i>Full sample</i>	<i>Passers</i>
	(1)	(2)
Delegated	0.002 (0.077)	−0.042 (0.078)
Bad outcome	0.556*** (0.139)	0.563*** (0.153)
Delegated × Bad outcome	−0.099 (0.095)	−0.072 (0.103)
Player B belief about Player A perf.	−0.083 (0.075)	−0.030 (0.082)
Age	−0.006 (0.013)	−0.013 (0.015)
Female	0.115 (0.272)	0.230 (0.291)
Socio-economic status	0.160 (0.126)	0.140 (0.129)
Went to uni	−0.827*** (0.320)	−0.706** (0.349)
Technology score	0.062 (0.118)	−0.023 (0.131)
Leadership position	0.062 (0.293)	−0.071 (0.321)
Constant	−0.382 (0.807)	−0.182 (0.864)
AME of Delegated (pp)	−1.8 ($p = 0.462$)	−2.9 ($p = 0.285$)
AME of Bad outcome (pp)	18.1*** ($p < 0.001$)	19.4*** ($p < 0.001$)
Observations	320	268
Subjects	80	67
Pseudo R^2	0.096	0.081

Note: Probit regression of an indicator for non-zero punishment ($1\{\text{punishment} > 0\}$) on a delegation indicator, a bad-outcome indicator, their interaction, Player B’s belief about Player A’s performance, and socio-demographic controls; the sample is stacked over the four (delegation, outcome) cells of the strategy method (4 observations per Player B). Column (1) uses the full Punishment-condition sample; Column (2) restricts to attention-check passers. Standard errors clustered at the Player B level. The *AME* rows report discrete average marginal effects on the probability of imposing any punishment, in percentage points (for Delegated, the change from $0 \rightarrow 1$ with the interaction moved consistently; for Bad outcome, likewise), computed by counterfactual prediction with p-values from a seeded 400-replication subject-cluster bootstrap. Significance: * $p < 0.10$; ** $p < 0.05$; *** $p < 0.01$ (two-sided).

Independence of beliefs from task performance. Figure 9 plots each of the four cell-level punishment beliefs and the two outcome-conditional belief differences against Player A’s task performance in the first ten rounds. Pearson correlations are reported in each panel; all six are small ($|r| \leq 0.21$) and statistically insignificant at conventional levels. In particular, the two belief differences that enter the regression in Table 3 are uncorrelated with task performance ($r = +0.12$, $p = 0.37$ for the good-outcome difference; $r = -0.15$, $p = 0.24$ for the bad-outcome difference). The within-subject correlation between beliefs and delegation reported in the main text therefore cannot be attributed to a shared dependence on task performance.

Figure 9: Player A punishment beliefs are independent of task performance

Player A punishment beliefs vs. task performance (Punishment condition, attention-check passers)



Note: Each panel plots one belief variable against Player A’s score in the first ten rounds, restricted to the *Punishment* condition and attention-check passers ($n = 61$). Points are jittered horizontally for legibility. Lines are OLS fits; annotations report Pearson r and its two-sided p -value. Top row: good-outcome beliefs (delegation, self, and the difference *self* – *delegation*). Bottom row: low-outcome beliefs.


B Additional figures

B.1 Hypothetical punishment

C Bonus computation

Player A spent on average 12 minutes on the experiment and earned an average bonus of £2.49 on top of the Prolific show-up fee; Player B spent on average 8 minutes and earned an average bonus of £2.30. Bonuses are computed deterministically from the cleaned data per the experimental rules: random-round task payoff (expectation over the uniform draw across

Figure 10: Hypothetical punishment by delegation and outcome scenario (*No-Punishment* condition)



Note: As in Figure 4, but for the *No-Punishment* condition (full sample $n = 81$, passers $n = 68$), in which Player Bs answered the same punishment questions but could not impose punishment in practice. For each (delegation, outcome) cell, the figure reports the average hypothetical punishment (left y -axis, in pounds; error bars are standard errors of the mean) and the share of Player Bs choosing non-zero punishment (right y -axis, in percent); the full sample is shown at full opacity and the attention-check passers overlaid at lower opacity. The left-axis scale matches Figure 4.

rounds 1–10); belief-about-own-score payoff (£0.30 times the probability mass placed on the realized score); belief-about-distribution payoff (expectation over the uniform random draw of which score is checked, paying £0.30 if the estimate is within 5 percentage points of the true population share); and—for Player A in the *Punishment* condition—expected punishment from the matched Player B’s strategy-method choice in the realized cell, with £0.10 per cell of strategy-method punishment cost for Player B. The reported bonuses exclude the risk-elicitation MPL payoff (the raw-data column was not retained in the cleaned files); including it would add roughly £0.30–£0.50 per subject.

D Ex-ante power calculations

The cell size of 80 subjects per role per condition was set to detect, at conventional power, the effect sizes anticipated from the responsibility-avoidance literature. We report the minimum detectable effect (MDE) for each preregistered test below, computed at $\alpha = 0.05$ (two-sided) and power = 0.80. For H1, the relevant test is a two-sample z -test of proportions (delegation rate in *Punishment* vs *No-Punishment*); the MDE is approximately invariant to the assumed baseline rate. For H2, the relevant test is a paired-mean-difference test on Player B’s punishment in the bad-outcome cell; because the within-subject standard deviation σ_d of the punishment difference is unknown ex ante, we report the MDE for a range of values of σ_d expressed as multiples of the equivalence bound used in Result 2 of the body of the paper ($\delta_a = £0.20$, equal to 10% of the maximum punishment of £2).

Table 9: Minimum detectable effects under the preregistered design

Hypothesis	Assumption	Minimum detectable effect
H1 (delegation, two-proportion test)	Baseline $p_0 = 0.30$	$\Delta = 0.216$ (21.6 pp)
H1 (delegation, two-proportion test)	Baseline $p_0 = 0.50$	$\Delta = 0.215$ (21.5 pp)
H1 (delegation, two-proportion test)	Baseline $p_0 = 0.70$	$\Delta = 0.179$ (17.9 pp)
H2 (paired, within-subject)	$\sigma_d = 0.5 \delta_a = \pounds 0.10$	$\Delta = \pounds 0.031$
H2 (paired, within-subject)	$\sigma_d = 1 \delta_a = \pounds 0.20$	$\Delta = \pounds 0.063$
H2 (paired, within-subject)	$\sigma_d = 2 \delta_a = \pounds 0.40$	$\Delta = \pounds 0.125$
H2 (paired, within-subject)	$\sigma_d = 3 \delta_a = \pounds 0.60$	$\Delta = \pounds 0.188$
H2 (paired, within-subject)	$\sigma_d = 4 \delta_a = \pounds 0.80$	$\Delta = \pounds 0.251$

Note: Both tests assume $n = 80$ per role per condition, $\alpha = 0.05$ two-sided, power = 0.80. For H2, σ_d is the within-subject standard deviation of the punishment difference (delegated minus non-delegated, conditional on bad outcome) and is expressed as a multiple of the anchor effect $\delta_a = \pounds 0.20$ (the equivalence bound used in Result 2 in the body of the paper, equal to 10% of the maximum punishment of $\pounds 2$). Δ is the corresponding minimum detectable effect in pounds.

E Preregistration and deviations

The experiment was preregistered at AsPredicted #133980, available at aspredicted.org/hs9az8.pdf and registered on 30 May 2023 prior to the main wave of data collection. The preregistration commits to the two research questions stated in the body as H1 and H2, the sample-size target (320 subjects total, 80 per role per condition), the desktop-only restriction, the use of the strategy method to elicit Player B’s punishment, and a set of primary statistical tests (parametric t -tests and non-parametric MWU and Kolmogorov–Smirnov tests for both comparisons; probit regressions for the delegation decision and linear regressions for punishment behavior with treatment, beliefs, and socio-demographics as predictors).

The preregistration also notes that a pilot of 25 observations had been collected in the *Punishment* treatment prior to registration; these observations are pooled with the main wave in all analyses, as preregistered. Two secondary analyses are pre-stated in the registry: the role of own-performance and beliefs about own performance (developed as M2 in Section 4.1), and the consistency of Player A’s beliefs about expected punishment with her delegation decision (developed as M3 in Section 4.1).

The table below summarizes every analysis reported in the body of the paper, the corresponding preregistered specification, and any deviation.

The full set of preregistered tests, the deviations from each, and the rationale appears in Appendix E


F Flow of participants

The flow of participants through the experiment, by condition and role, is shown in Figure 11. Recruitment-funnel counts (invited / consented) are pending the Prolific audit log.

Table 10: Preregistered tests, reported tests, and deviations

Preregistered specification	As reported	Deviation / rationale
H1: between-treatment test of delegation rates (<i>Punishment</i> vs <i>No-Punishment</i>); preregistered tests are parametric <i>t</i> -test, MWU, KS. Direction not specified in the registry.	Pearson Chi-squared and Fisher’s exact, two-sided ($p = 0.033$ / $p = 0.041$).	Test family substituted (Chi-squared and Fisher’s exact are the standard tests for a binary outcome; the preregistered <i>t</i> -test, MWU and KS tests on a binary outcome are equivalent to the proportion test). The realized direction is opposite to that anticipated by the responsibility-avoidance literature on which we drew.
H2: paired test of Player B punishment in the bad-outcome cell, delegated vs non-delegated, attention-check passers only	Paired <i>t</i> -test ($p = 0.153$); Wilcoxon signed-rank ($p = 0.533$).	As preregistered.
Probit / logit of delegation on condition + controls (Table 1, Cols (1)–(2))	Logit, as reported.	As preregistered. (Probit and logit are equivalent for our purposes; we report logit for ease of marginal-effect interpretation.)
Logit on <i>Punishment</i> subsample with belief differences (Col (3))	As reported.	Specification preregistered; sample restriction follows the preregistered attention-check rule.
Pre-stated secondary analysis: own-performance and beliefs about own performance	Reported as M2 in Section 4.1.	As preregistered (secondary).
Pre-stated secondary analysis: beliefs about expected punishment vs delegation decision	Reported as M3 in Section 4.1.	As preregistered (secondary).
Mechanism analyses M1 (process ownership) and M4 (experimenter demand) (Section 4.1)	As reported.	<i>Not preregistered.</i> Treated as exploratory.

Figure 11: Flow of participants through the experiment



A flow diagram showing recruitment, attrition, attention-check exclusion, and the analysis sample, broken down by condition and role, appears in Appendix F; in particular, attrition rates do not differ meaningfully across the *Punishment* and *No-Punishment* conditions.

We preregistered exclusion of subjects who failed the attention check on the screen at which their main outcome variable was elicited. For Player A, this is the belief-elicitation screen; for Player B, the punishment-elicitation screen. Of 161 Player Bs, 135 pass the attention check (83.9%); the corresponding pass rate for Player A is 131/161 (81.4%). All main results are robust to including subjects who fail the attention check; we report these robustness checks in Appendix A.2.


FB: maybe move to first set of results. Also the appendix does not show anything and the default should be including everyone.

G Additional Figures

G.1 Experimental timelines


G.2 Punishment frequencies

Figure 12: Experimental timeline — Player A




Note: Screen sequence presented to Player A. The shaded screen carries the attention check used to identify subjects excluded from belief-related analyses. The final-prediction-and-confidence screen is skipped if Player A delegates to the algorithm.

Figure 13: Experimental timeline — Player B



Note: Screen sequence presented to Player B. The shaded screen carries the attention check used to identify subjects excluded from punishment-related analyses. In the *No-Punishment* condition, no actual punishment is implemented; the elicitation is hypothetical and serves to maintain comparability of the screen sequence across conditions.

Figure 14: Frequencies of punishment and beliefs about punishment



Note: Share of Player Bs choosing strictly positive punishment in each of the four scenarios (solid bars) and share of Player As assigning a strictly positive expected punishment in each scenario (hollow bars).